

Intro. to Computer Architecture

Daniel Page

Department of Computer Science,
University Of Bristol,
Merchant Venturers Building,
Woodland Road,
Bristol, BS8 1UB. UK.
(csdsp@bristol.ac.uk)

January 9, 2018

Keep in mind there are *two* PDFs available (of which this is the latter):

1. a PDF of examinable material used as lecture slides, and
2. a PDF of non-examinable, extra material:
 - ▶ the associated notes page may be pre-populated with extra, written explanation of material covered in lecture(s), plus
 - ▶ anything with a “grey’ed out” header/footer represents extra material which is useful and/or interesting but out of scope (and hence not covered).

Notes:

Notes:

► Agenda:

1. comments, questions, recap, then
2. a short overview of (hardware-oriented) strategies for **testing** and **debugging**.

Notes:

Quote

Beware of bugs in the above code; I have only proved it correct, not tried it.

– Knuth (http://en.wikiquote.org/wiki/Donald_Knuth)

Quote

Everyone knows that debugging is twice as hard as writing a program in the first place. So if you're as clever as you can be when you write it, how will you ever debug it?

– Kernighan (http://en.wikiquote.org/wiki/Brian_Kernighan)

Notes:

► Example:

- The Intel Pentium (P5) model processor has a 273-pin package; in the worse case there are

$$2^{273} = 151771007205135083665582961470587414581438 \backslash 03430094840009779784451085189728165691392$$

possible inputs (ignoring outputs, and maintained state).

- In 1994, a researcher discovered the processor had a bug:

- The processor should have been able to compute

$$(4195835 \cdot 3145727) / 3145727 = 4195835.$$

- The so-called FDIV bug [5] resulted in

$$(4195835 \cdot 3145727) / 3145727 = 4195579.$$

- Intel recalled the defective processors, at an estimated cost of \$475 million.

Notes:

Definition

Given

specification \leadsto design \leadsto implementation \leadsto **Design Under Test (DUT)**

we should carefully distinguish between

validation	\Rightarrow	check whether specification matches requirements (i.e., is fit for purpose)
	\approx	“is this the correct specification?”
verification	\Rightarrow	check (or <i>prove</i>) whether implementation matches specification
	\approx	“is the implementation correct?”
testing	\Rightarrow	check (or <i>prove</i>) whether DUT matches specification
	\approx	“is the DUT correct?”

Notes:

Definition

A **fault** describes a (static) bug in the DUT:

- ▶ **design faults** (or **functional faults**) i.e., design is incorrect in some way,
- ▶ **implementation faults**, e.g., design is correct but synthesis produced an incorrect implementation,
- ▶ **manufacture faults**, i.e., implementation is correct but manufacturing defects produce an incorrect DUT,
- ▶ **operational faults**, e.g., due to wear (or “old age”) or the impact of operational limits.

Definition

An **error** is an instance of the fault being triggered during (dynamic) use of the DUT.

Definition

A **failure** is the *effect* of an error(s); perhaps the DUT malfunctions, for example, or perhaps it explodes.

Notes:

- Another important classification wrt. faults is whether they are
 - **permanent**, or
 - **transient**,


in either case noting a fault may only occur after a **burn-in** period: this motivates explicit burn-in tests [2], which could be viewed as a type of stress test.

- Example faults include the following:
 - timing (or delay) faults, where a signal may settle to the correct value but more slowly or quickly than expected,
 - stuck-at faults, where a signal is permanently 0 or 1,
 - closed connection (or bridged connection) faults, where two wires are “fused” when they should be separate,
 - open connection faults, where a wire is “broken” so that one end is disconnected from the other.

A case-study: testing a ripple-carry adder (1)

- ▶ **Recap**: we produced the following design for (ripple-carry) addition

Circuit



- ▶ **Problem**: if we implement this design, how do we
 - ▶ test whether the implementation is correct, *and*
 - ▶ diagnose (or debug) the implementation (and/or design) if not.
- ▶ **Solution**: we need (at least) two components ...

Notes:

A case-study: testing a ripple-carry adder (2)

Definition

The term **test fixture** is typically used to capture the *physical* infrastructure for testing, e.g., the equipment used to interact with a DUT and thereby provide input (or stimuli) and inspect output.

Definition

The term **test bench** is typically used to capture the *logical* infrastructure for testing, e.g.,

1. a set of inputs,
2. a mechanism to produce an expected output given an input,
3. a mechanism to check the actual vs. expected output.

Note that where the DUT is simulated, the test bench can be entirely virtual; otherwise, a test fixture is required to interact with the physical DUT.



Notes:

A case-study: testing a ripple-carry adder (3)

Component #1: a test fixture

- **Problem:** how do form connectivity with the DUT?
- **Solution:** given a packaged IC, or PCB, we actually have (at least) two options
 1. **in-circuit test**, *with* test fixture ⇒ bed of nails
 2. **in-circuit test**, *without* test fixture ⇒ flying lead

e.g.,



both of which are easier where **test points** are available.

Notes:

A case-study: testing a ripple-carry adder (4)

Component #1: a test fixture

► **Problem:** how do we interact with (e.g., control) the DUT?


► **Solution:** **boundary scan** [1].

► **Joint Test Action Group (JTAG) [3]** is a common approach.

► The basic idea is to include two on-chip components

1. a **scan-chain** that links together all the latches (resp. flip-flops), and
2. a **JTAG interface** that allows external access to the scan-chain,

i.e.,



1. Test Data In (TDI),
2. Test Data Out (TDO),
3. Test Clock (TCK),
4. Test Mode Select (TMS), and
5. Test Reset (TRST) which is optional.

► We then use these components to interact with the DUT ...

Notes:

A case-study: testing a ripple-carry adder (4)

Component #1: a test fixture

► **Problem:** how do we interact with (e.g., control) the DUT?


► **Solution:** **boundary scan** [1].

► **Joint Test Action Group (JTAG) [3]** is a common approach.

► The basic idea is to include two on-chip components

1. a **scan-chain** that links together all the latches (resp. flip-flops), and
2. a **JTAG interface** that allows external access to the scan-chain,

i.e.,




1. Test Data In (TDI),
2. Test Data Out (TDO),
3. Test Clock (TCK),
4. Test Mode Select (TMS), and
5. Test Reset (TRST) which is optional.

► We then use these components to interact with the DUT ...

Notes:


Example



1. Initially, the circuit computes results as normal.

Notes:

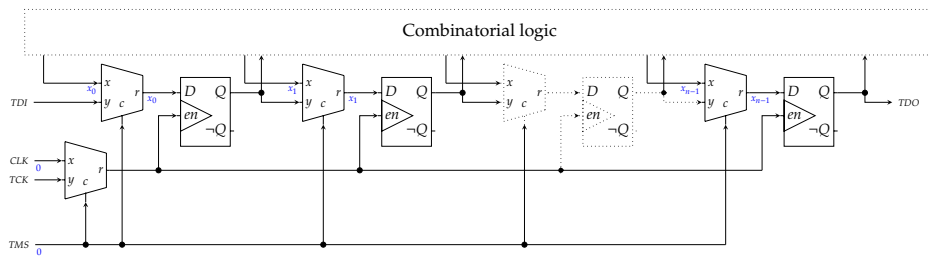
Example



1. Initially, the circuit computes results as normal.

Notes:

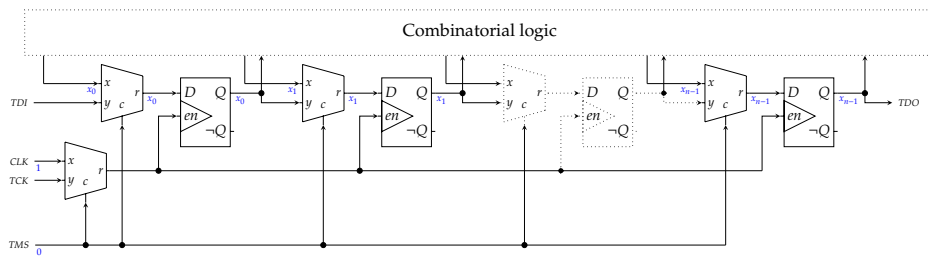
Example



1. Initially, the circuit computes results as normal.

Notes:

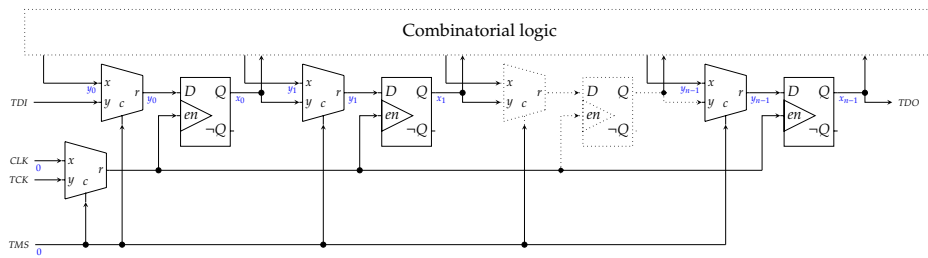
Example



1. Initially, the circuit computes results as normal.

Notes:

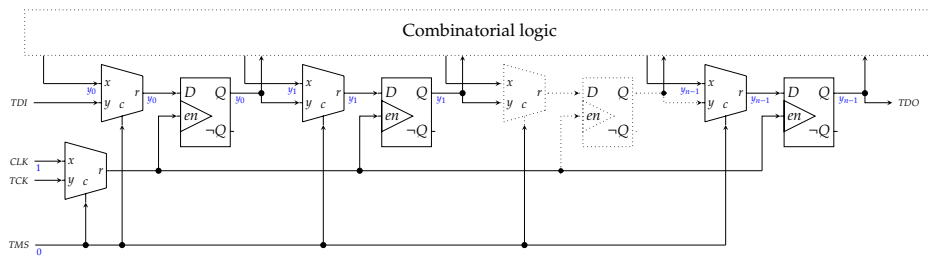
Example



1. Initially, the circuit computes results as normal.

Notes:

Example

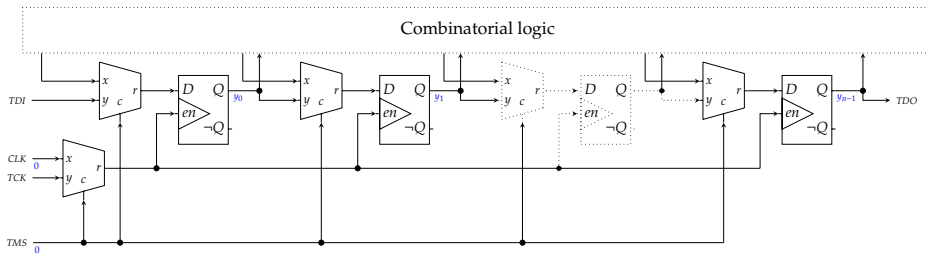


1. Initially, the circuit computes results as normal.

Notes:

A case-study: testing a ripple-carry adder (5)

Example

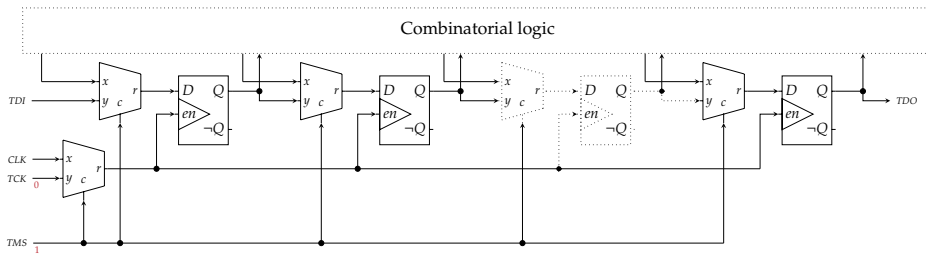


- Initially, the circuit computes results as normal.

Notes:

A case-study: testing a ripple-carry adder (5)

Example

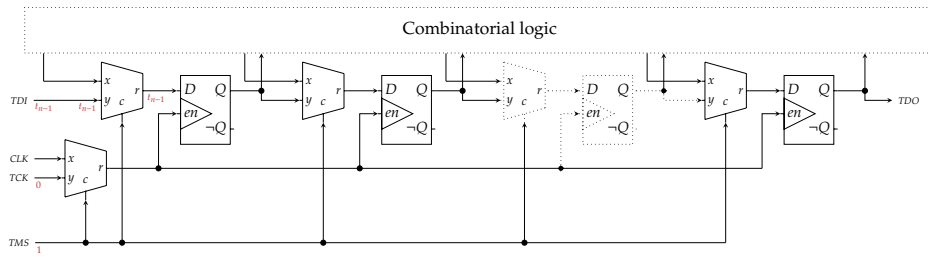


1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.

Notes:

Component #1: a test fixture

Example



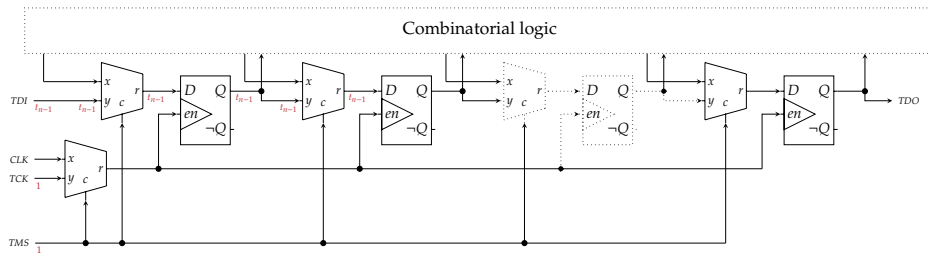
- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:

A case-study: testing a ripple-carry adder (5)

Component #1: a test fixture

Example

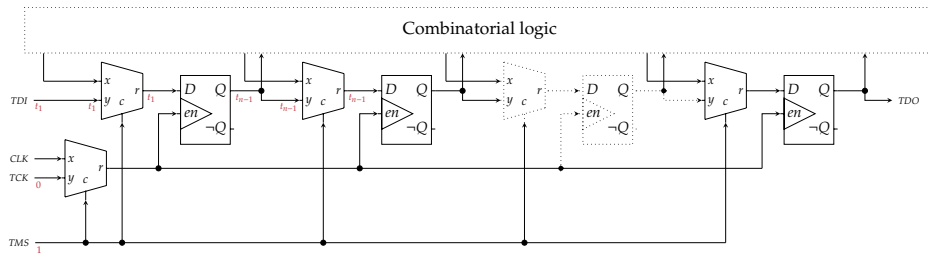


- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:

Component #1: a test fixture

Example



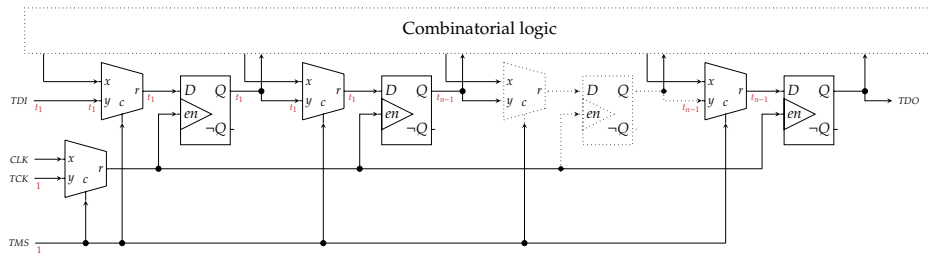
- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:

A case-study: testing a ripple-carry adder (5)

Component #1: a test fixture

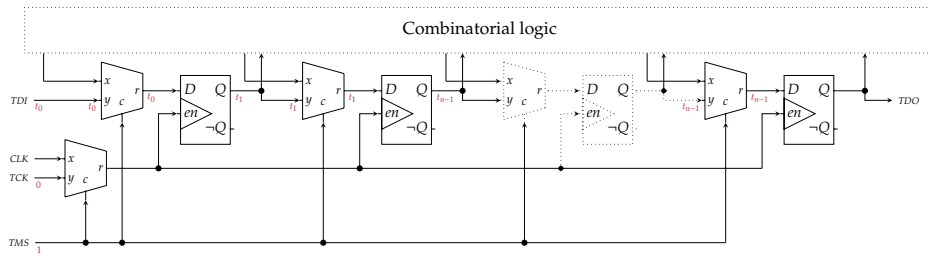
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:


Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:

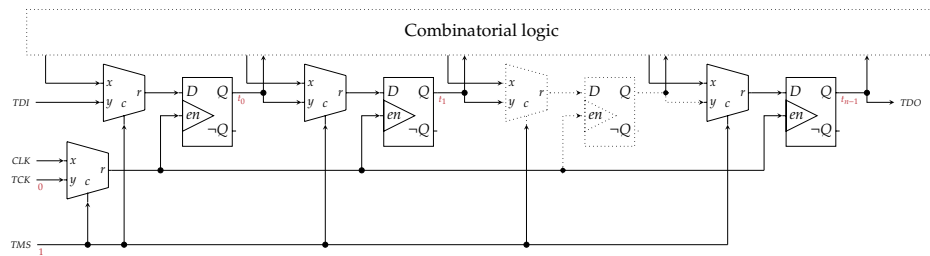
Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:

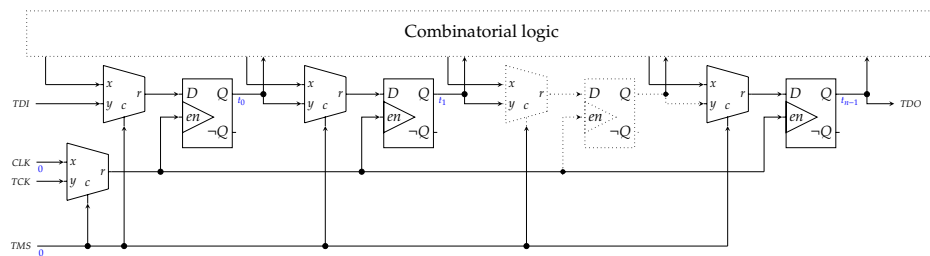
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.

Notes:

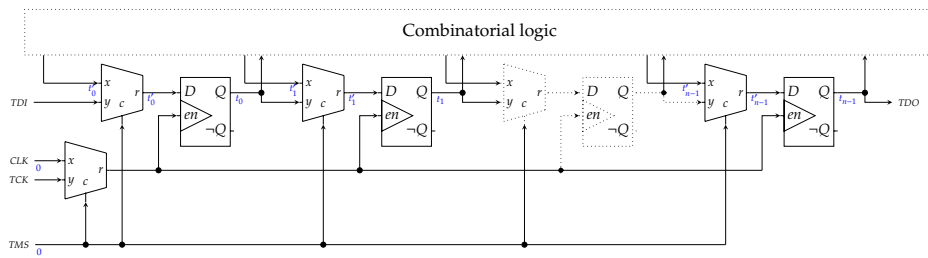
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
- The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.

Notes:

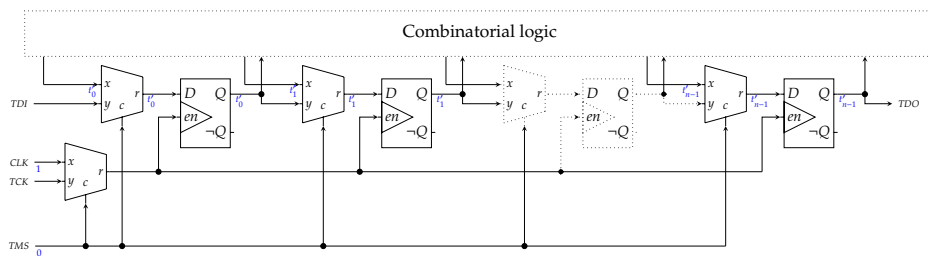
Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
4. The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.

Notes:

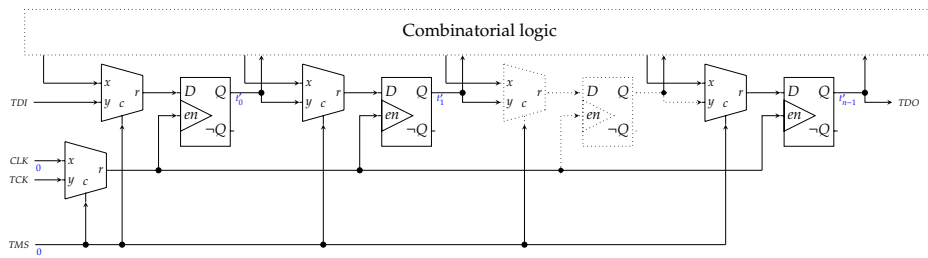
Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
4. The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.

Notes:

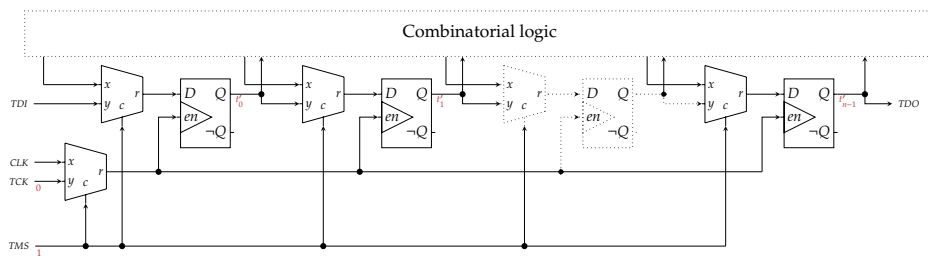
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
- The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.

Notes:

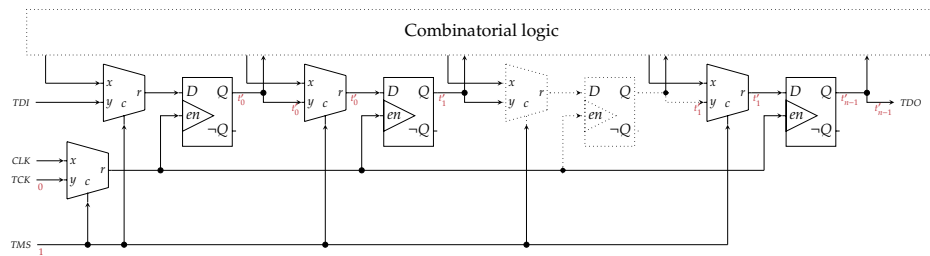
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
- The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
- The circuit is placed in **debug mode** using TMS.

Notes:

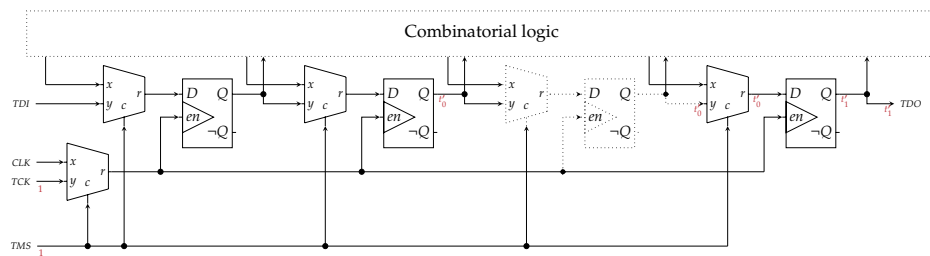
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
- The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
- The circuit is placed in **debug mode** using TMS.
- The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:

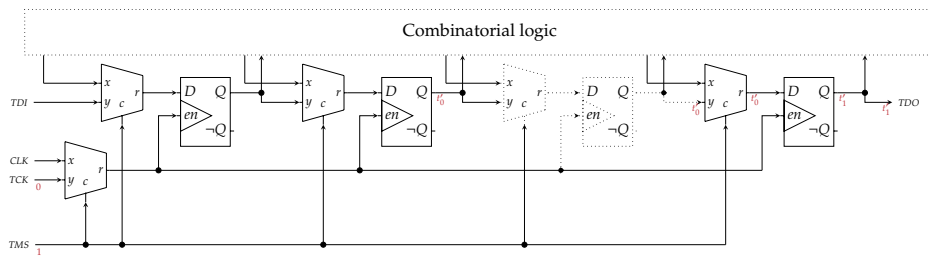
Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
- The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
- The circuit is placed in **debug mode** using TMS.
- The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:

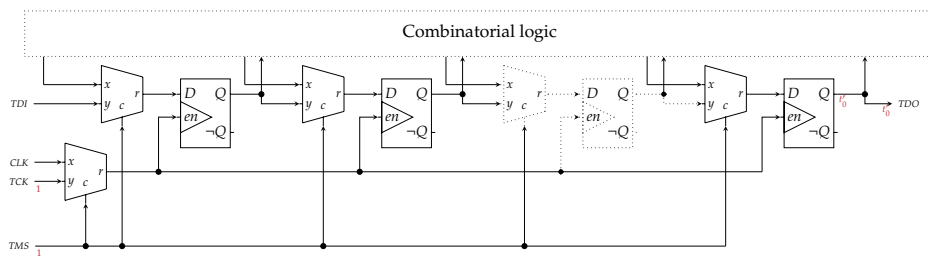
Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
4. The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
5. The circuit is placed in **debug mode** using TMS.
6. The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:

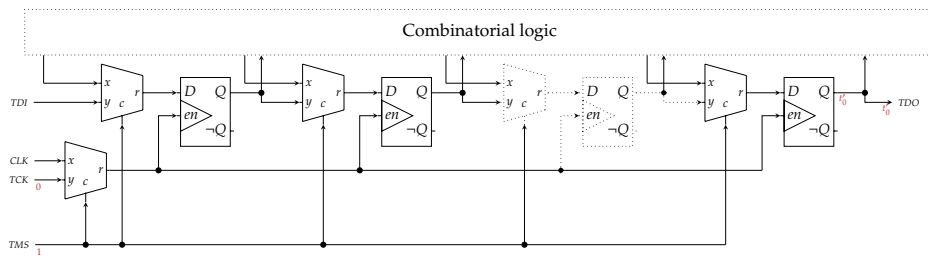
Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
4. The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
5. The circuit is placed in **debug mode** using TMS.
6. The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:

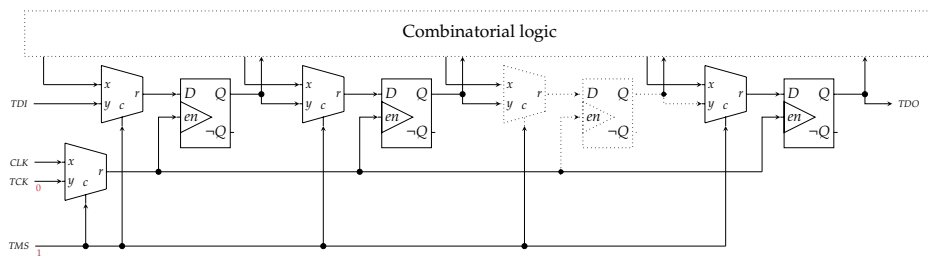
Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
4. The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
5. The circuit is placed in **debug mode** using TMS.
6. The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:


Example



1. Initially, the circuit computes results as normal.
2. Then, at some point, it is placed in **debug mode** using TMS.
3. The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
4. The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
5. The circuit is placed in **debug mode** using TMS.
6. The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:

Example



- Initially, the circuit computes results as normal.
- Then, at some point, it is placed in **debug mode** using TMS.
- The test vector is scanned in through TDI; each edge on TCK reads from TDI and shifts the scan-chain along.
- The circuit is placed in **normal mode** using TMS and clocked until the required result is computed.
- The circuit is placed in **debug mode** using TMS.
- The result vector is scanned out through TDI; each edge on TCK writes to TDO and shifts the scan-chain along.

Notes:

Definition

A **test vector** (or **test pattern**) is a single input and associated output; using a *set* of n test vectors

- increases the level of coverage (i.e., proportion of DUT exercised), and
- increases our *confidence* the DUT is correct.

► **Item:** given an oracle [4] O for functionality f as realised by the DUT,

- construct a set

$$\{(x_0, r_0 = O(x_0)), (x_1, r_1 = O(x_1)), \dots, (x_{n-1}, r_{n-1} = O(x_{n-1}))\}$$

of test vectors, then

- interact with the DUT to check

$$f(x_i) \stackrel{?}{=} O(x_i)$$

for all $0 < i < n$.

Notes:

Definition

A (set of) test vector(s) could be classified at least per

1. **exhaustive** :
2. **heuristic** :
3. **diagnostic** :

Definition

The generation and use of a (set of) test vector(s) forms a wider taxonomy still: they could be deemed

1. **online** or **offline**,
2. **static** or **adaptive**, and
3. **random** or **deterministic**.

Notes:

Definition

A (set of) test vector(s) could be classified at least per

1. **exhaustive** : $x \in \{-2^{n-1}, \dots, 2^{n-1} - 1\}$
2. **heuristic** : $\begin{cases} \text{small cases st. } x \in \{\pm 0, \pm 1, \dots, \pm 2^m - 1\} \text{ for } m \ll n \\ \text{boundary cases st. } x \in \{-2^{n-1}, -1, 0, +1, 2^{n-1} - 1\} \end{cases}$
3. **diagnostic** : an x selected st. if $f(x)$ fails, we infer that a specific fault exists

Definition

The generation and use of a (set of) test vector(s) forms a wider taxonomy still: they could be deemed


1. **online** or **offline**,
2. **static** or **adaptive**, and
3. **random** or **deterministic**.

Notes:

A case-study: testing a ripple-carry adder (8)

Component #2: a test bench

- **Problem:** how could we organise the test bench?
- **Solution(s):**




1. stand-alone test bench \approx external test bench

A case-study: testing a ripple-carry adder (8)

Component #2: a test bench

- **Problem:** how could we organise the test bench?
- **Solution(s):**



1. stand-alone test bench \approx external test bench

Notes:

- It is obviously possible to enumerate a list of positive and negative features for each approach; one or other may be the better option for a given context. Although non-exhaustive, the following attempts to capture some of the main points:
 1. The stand-alone test bench:
 - +ve: can pre-compute test vectors, using an oracle that might be represented by a database, or an alternative implementation, for example,
 - +ve: test plan can be adaptive,
 - ve: tests are typically applied after manufacture only, so not close to time of use,
 - ve: is less resource efficient in the sense it does require a dedicated test fixture,
 - +ve: is more resource efficient in the sense there is 1 test fixture for n DUTs.
 2. The BIST approach:
 - ve: cannot (as easily) pre-compute test vectors, since the storage requirement is typically too great,
 - ve: cannot cope (as easily) with faults that occur in the test mechanism itself,
 - +ve: tests can be applied close to time of use,
 - ve: test plan cannot (as easily) be adaptive, in the sense it must be fixed when the DUT is manufactured,
 - +ve: is more resource efficient in the sense it does not require a dedicated test fixture,
 - ve: is less resource efficient in the sense there are n test benches and test fixtures for n DUTs (because they are effectively built into *each* DUT).
- The concept of *arithmetic* BIST, for example relies on testing arithmetic identities such as

$$(x + y) - (x + y) \stackrel{?}{=} 0$$

or

$$(x + y) + z \stackrel{?}{=} x + (y + z).$$

where the advantage of doing this is we can just select random x , y and z and check whether the comparison holds (vs. having to construct an explicit set of test vectors). Clearly this concept translates to non-arithmetic cases *if* similar identities are possible.

Notes:

- It is obviously possible to enumerate a list of positive and negative features for each approach; one or other may be the better option for a given context. Although non-exhaustive, the following attempts to capture some of the main points:
 1. The stand-alone test bench:
 - +ve: can pre-compute test vectors, using an oracle that might be represented by a database, or an alternative implementation, for example,
 - +ve: test plan can be adaptive,
 - ve: tests are typically applied after manufacture only, so not close to time of use,
 - ve: is less resource efficient in the sense it does require a dedicated test fixture,
 - +ve: is more resource efficient in the sense there is 1 test fixture for n DUTs.
 2. The BIST approach:
 - ve: cannot (as easily) pre-compute test vectors, since the storage requirement is typically too great,
 - ve: cannot cope (as easily) with faults that occur in the test mechanism itself,
 - +ve: tests can be applied close to time of use,
 - ve: test plan cannot (as easily) be adaptive, in the sense it must be fixed when the DUT is manufactured,
 - +ve: is more resource efficient in the sense it does not require a dedicated test fixture,
 - ve: is less resource efficient in the sense there are n test benches and test fixtures for n DUTs (because they are effectively built into *each* DUT).
- The concept of *arithmetic* BIST, for example relies on testing arithmetic identities such as


$$(x + y) - (x + y) \stackrel{?}{=} 0$$

or

$$(x + y) + z \stackrel{?}{=} x + (y + z).$$

where the advantage of doing this is we can just select random x , y and z and check whether the comparison holds (vs. having to construct an explicit set of test vectors). Clearly this concept translates to non-arithmetic cases *if* similar identities are possible.

- **Problem:** how could we organise the test bench?
- **Solution(s):**



2. Built-In Self-Test (BIST) \simeq *internal* test bench

Conclusions

► Take away points:

1. On the whole, similar best-practices can be leveraged in software and hardware testing ...
2. ... but a bug in software means re-compilation whereas a bug in hardware means re-manufacturing, so, in a sense

verification \gg testing.

3. Robust methodology and a **design-for-test** ethos ideally form first class requirements:

- formulate a **test plan**, and
- consider infrastructure for said plan,

at the *design* stage, to then make it easier and/or more useful when applied to a DUT.

4. This general area is often viewed as secondary to “getting the job done” (i.e., developing a design); the *costs* involved mean it’s a larger element of “the job” than you might expect.

Notes:

- It is obviously possible to enumerate a list of positive and negative features for each approach; one or other may be the better option for a given context. Although non-exhaustive, the following attempts to capture some of the main points:

1. The stand-alone test bench:

- +ve: can pre-compute test vectors, using an oracle that might be represented by a database, or an alternative implementation, for example,
- +ve: test plan can be adaptive,
- ve: tests are typically applied after manufacture only, so not close to time of use,
- ve: is less resource efficient in the sense it does require a dedicated test fixture,
- +ve: is more resource efficient in the sense there is 1 test fixture for n DUTs.

2. The BIST approach:

- ve: cannot (as easily) pre-compute test vectors, since the storage requirement is typically too great,
- ve: cannot cope (as easily) with faults that occur in the test mechanism itself,
- +ve: tests can be applied close to time of use,
- ve: test plan cannot (as easily) be adaptive, in the sense it must be fixed when the DUT is manufactured,
- +ve: is more resource efficient in the sense it does not require a dedicated test fixture,
- ve: is less resource efficient in the sense there are n test benches and test fixtures for n DUTs (because they are effectively built into *each* DUT).

- The concept of *arithmetic* BIST, for example relies on testing arithmetic identities such as

$$(x + y) - (x + y) \stackrel{?}{=} 0$$

or

$$(x + y) + z \stackrel{?}{=} x + (y + z).$$

where the advantage of doing this is we can just select random x , y and z and check whether the comparison holds (vs. having to construct an explicit set of test vectors). Clearly this concept translates to non-arithmetic cases *if* similar identities are possible.

Notes:

References

- [1] *Wikipedia: Boundary scan*. URL: http://en.wikipedia.org/wiki/Boundary_scan (see pp. 21, 23).
- [2] *Wikipedia: Burn-in*. URL: <http://en.wikipedia.org/wiki/Burn-in> (see p. 14).
- [3] *Wikipedia: Joint Test Action Group (JTAG)*. URL: http://en.wikipedia.org/wiki/Joint_Test_Action_Group (see pp. 21, 23).
- [4] *Wikipedia: Oracle*. URL: [http://en.wikipedia.org/wiki/Oracle_\(software_testing\)](http://en.wikipedia.org/wiki/Oracle_(software_testing)) (see p. 79).
- [5] *Wikipedia: Pentium FDIV bug*. URL: http://en.wikipedia.org/wiki/Pentium_FDIV_bug (see p. 9).

Notes: