

Intro. to Computer Architecture

Daniel Page

Department of Computer Science,
University Of Bristol,
Merchant Venturers Building,
Woodland Road,
Bristol, BS8 1UB. UK.
(csdsp@bristol.ac.uk)

November 3, 2017

Keep in mind there are *two* PDFs available (of which this is the latter):

1. a PDF of examinable material used as lecture slides, and
2. a PDF of non-examinable, extra material:
 - ▶ the associated notes page may be pre-populated with extra, written explanation of material covered in lecture(s), plus
 - ▶ anything with a “grey’ed out” header/footer represents extra material which is useful and/or interesting but out of scope (and hence not covered).

Notes:

Notes:

- **Agenda:**
- 1. comments, questions, recap, then
 - 2. a general-purpose implementation technique and related trade-offs, namely
 - the concept of **pipelining**, and
 - trade-offs between area, latency, and throughput.

Notes:


Sequential use of combinatorial stages \simeq pipelining (1)
An analogy: production lines

Worker #1	Worker #2
Worker #3	Worker #4

Notes:

Sequential use of combinatorial stages \approx pipelining (1)

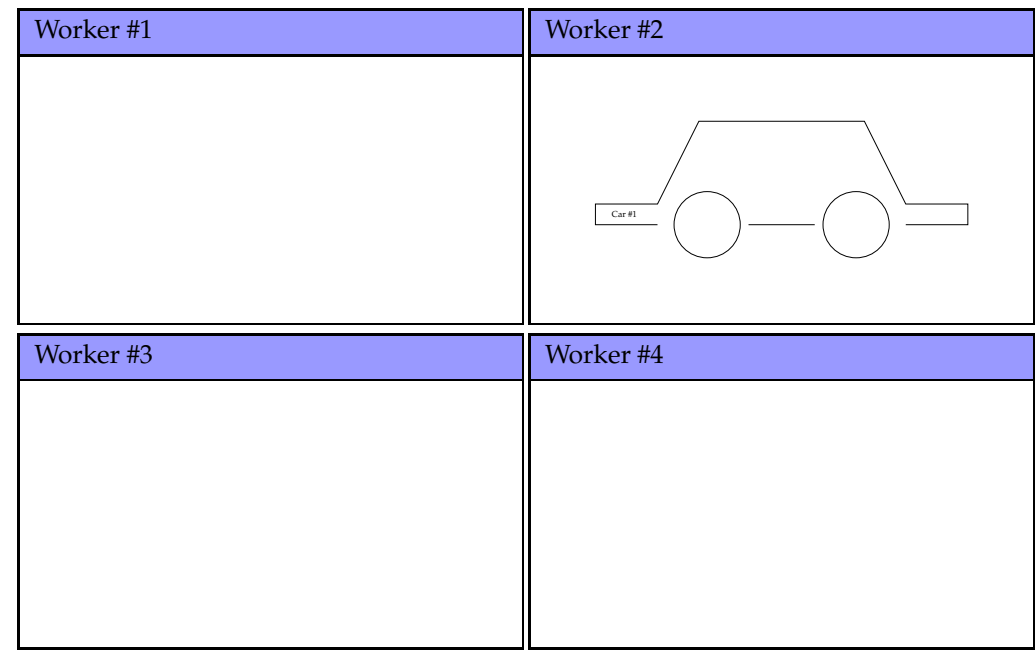
An analogy: production lines



Notes:


Sequential use of combinatorial stages \approx pipelining (1)

An analogy: production lines



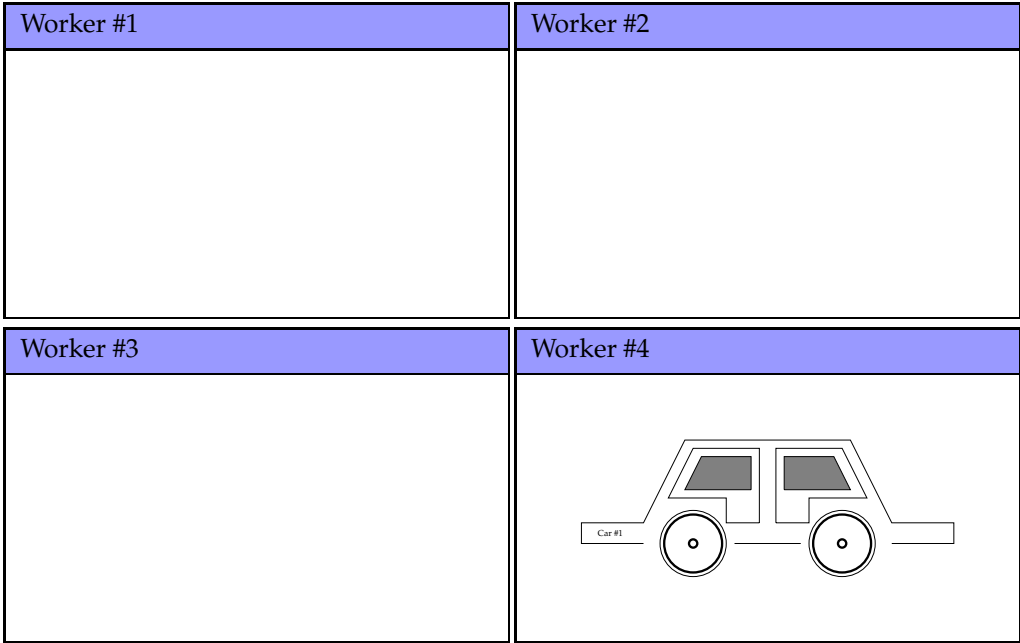
Notes:

Sequential use of combinatorial stages \approx pipelining (1)
An analogy: production lines



Notes:

Sequential use of combinatorial stages \approx pipelining (1)
An analogy: production lines




Notes:

Sequential use of combinatorial stages \approx pipelining (1)
An analogy: production lines

Worker #1	Worker #2
Worker #3	Worker #4

Notes:


Sequential use of combinatorial stages \approx pipelining (1)
An analogy: production lines

Worker #1	Worker #2
	
Worker #3	Worker #4

Notes:

Sequential use of combinatorial stages \approx pipelining (1)

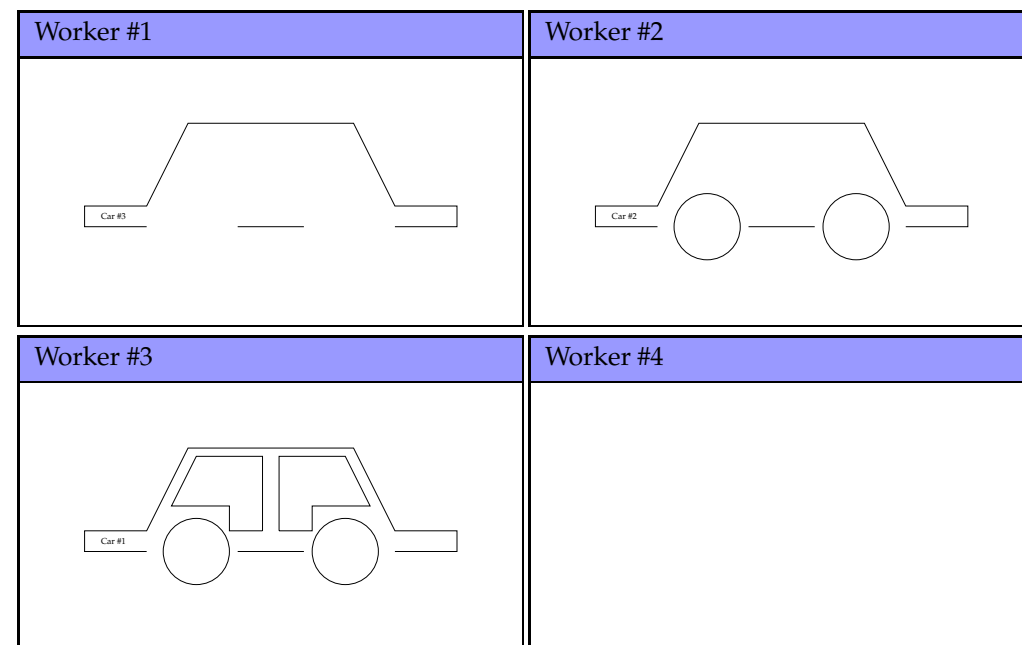
An analogy: production lines



Notes:

Sequential use of combinatorial stages \approx pipelining (1)


An analogy: production lines



Notes:

Sequential use of combinatorial stages \approx pipelining (1)

An analogy: production lines



Notes:

Sequential use of combinatorial stages \approx pipelining (2)

An analogy: production lines

Definition

The **latency** is the total time elapsed before a given input is operated on to produce an output; this is simply the sum of the latencies of each stage.

Definition

The **throughput** (or **bandwidth**) is the rate at which new inputs can be supplied (resp. outputs collected).

► Good news:

- Based on the metrics above we find

without pipeline \leadsto latency = 4, throughput = 1/4
with pipeline \leadsto latency = 4, throughput = 1

suggesting that an (idealised) n -stage yields an n -fold improvement in throughput.

- The production of one car is *still* sequential, but we used parallelism to *hide* latency wrt. a given car.

Notes:

Sequential use of combinatorial stages \approx pipelining (2)

An analogy: production lines

Definition

The **latency** is the total time elapsed before a given input is operated on to produce an output; this is simply the sum of the latencies of each stage.

Definition

The **throughput** (or **bandwidth**) is the rate at which new inputs can be supplied (resp. outputs collected).

► Bad news:

- If we can't **advance** the production line for some reason (e.g., a stage is delayed), we say it has **stalled**; this is bad, since it reduces utilisation.
- Advancement is limited by the slowest stage; to minimise idle time, balance is needed between the workload of stages.
- Even then, there is overhead associated with *all* stages (e.g., communicating between stages takes some time).

Notes:

Sequential use of combinatorial stages \approx pipelining (3)

Application to digital logic

► Idea: apply the production line approach to *circuits*, i.e.,

1. split some combinatorial logic X into a **pipeline** of n stages, say X_i for $0 \leq i < n$,
2. have each stage perform one step of the computation,
3. allow with multiple independent computations to overlap, with partial (or active) cases deemed **in-flight**.

Notes:

Sequential use of combinatorial stages \simeq pipelining (3)

Application to digital logic


- **Problem #1:** how to split X into stages, or rather
 1. *where* should we split it (which depends heavily on what X is), and
 2. *how* should we split it?
- **Problem #2:** how do we control the resulting pipeline?

Notes:

Sequential use of combinatorial stages \simeq pipelining (4)

Application to digital logic

- **Solution #1:** consider a (simplistic) example,



from which we can infer that


- the slowest stage dictates how fast we can advance the pipeline, so
- we need to balance the stages st. idleness is minimised (i.e., we avoid one stage waiting for another).

Notes:

Sequential use of combinatorial stages \simeq pipelining (4)

Application to digital logic

- **Solution #1:** consider a (simplistic) example,



from which we can infer that

- the slowest stage dictates how fast we can advance the pipeline, so
- we need to balance the stages st. idleness is minimised (i.e., we avoid one stage waiting for another).

Notes:

Sequential use of combinatorial stages \simeq pipelining (4)

Application to digital logic

- **Solution #1:** consider a (simplistic) example,



from which we can infer that


- the slowest stage dictates how fast we can advance the pipeline, so
- we need to balance the stages st. idleness is minimised (i.e., we avoid one stage waiting for another).

Notes:

Sequential use of combinatorial stages \simeq pipelining (4)

Application to digital logic

- **Solution #1:** consider a (simplistic) example,



from which we can infer that


- the slowest stage dictates how fast we can advance the pipeline, so
- we need to balance the stages st. idleness is minimised (i.e., we avoid one stage waiting for another).

Notes:

Sequential use of combinatorial stages \simeq pipelining (5)

Application to digital logic

- **Solution #2:**
 1. insert a **pipeline register** between adjacent stages, and
 2. advance the pipeline on each positive edge of some control signal,i.e.,



where

1. for a *synchronous* pipeline, *adv* is basically a global clock signal, but
2. for an *asynchronous* pipeline, some form of protocol between the stages controls *adv_i*: in a sense the stages decide when to advance when they're ready.

Notes:

- You should view the solid lines as permanent connections, and the dashed lines as connections that are used during update: the dashed lines hence isolate the stages from each other, so they only interact at the point when some R_{i+1} is updated to hold t_i .

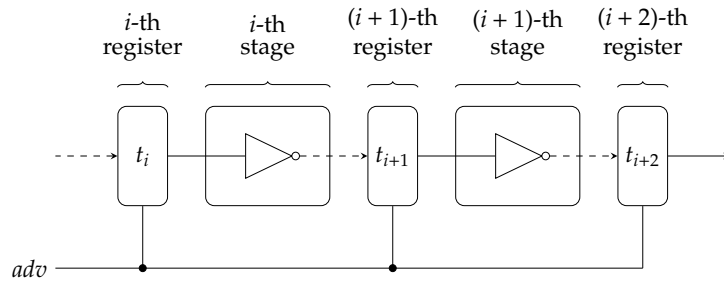
Sequential use of combinatorial stages \simeq pipelining (5)

Application to digital logic

► Solution #2:

1. insert a **pipeline register** between adjacent stages, and
2. advance the pipeline on each positive edge of some control signal,

i.e.,



where

1. for a *synchronous* pipeline, adv is basically a global clock signal, but
2. for an *asynchronous* pipeline, some form of protocol between the stages controls adv_i : in a sense the stages decide when to advance when they're ready.

Notes:

- You should view the solid lines as permanent connections, and the dashed lines as connections that are used during update: the dashed lines hence isolate the stages from each other, so they only interact at the point when some R_{i+1} is updated to hold t_i .


Sequential use of combinatorial stages \simeq pipelining (5)

Application to digital logic

► Solution #2:

1. insert a **pipeline register** between adjacent stages, and
2. advance the pipeline on each positive edge of some control signal,

i.e.,



where

1. for a *synchronous* pipeline, adv is basically a global clock signal, but
2. for an *asynchronous* pipeline, some form of protocol between the stages controls adv_i : in a sense the stages decide when to advance when they're ready.

Notes:

- You should view the solid lines as permanent connections, and the dashed lines as connections that are used during update: the dashed lines hence isolate the stages from each other, so they only interact at the point when some R_{i+1} is updated to hold t_i .

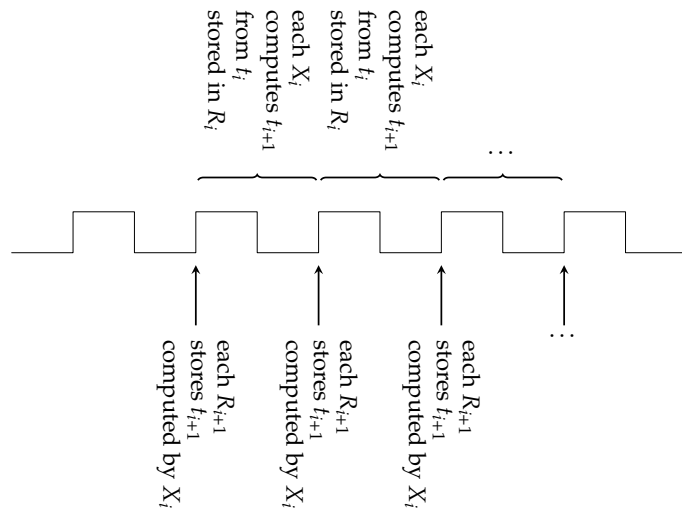
- ▶ *Why?!*
 - ▶ Imagine in the j -th clock cycle the i -th stage X_i computes partial result t_i required by the $(i + 1)$ -th stage.
 - ▶ If the stages are connected directly via a wire,
 1. when the i -th stage changes t_i , this disrupts computation by the $(i + 1)$ -th stage,
 2. the i -th and $(i + 1)$ -th stages cannot perform different in-flight computations, because there is nowhere to store the associated (partial) inputs.
- ▶ So,
 - +ve: If the stages are connected indirectly via a pipeline register, say R_i , the $(i + 1)$ -th stage can have a separate, stable input.
 - ve: Each pipeline register takes time to operate, so adds to the total latency: latency and throughput become a trade-off.

Notes:

Sequential use of combinatorial stages \approx pipelining (7)

Application to digital logic

Example




Notes:

Sequential use of combinatorial stages \approx pipelining (8)

Example: a generic component

Example

Consider some generic, combinatorial logic X :



1. For the standard design

- ▶ the latency is $300 + 20\text{ns} = 320\text{ns}$, while
- ▶ the throughput is $1/320\text{ns} = 3.12 \times 10^6 \text{ operations/s}$

if we measure the latency of computing X and storing the result.


Notes:

Sequential use of combinatorial stages \approx pipelining (8)

Example: a generic component

Example

Consider some generic, combinatorial logic X :



2. But, for a 3-stage pipeline (using the same measure)


- ▶ the latency is $100 + 20 + 100 + 20 + 100 + 20\text{ns} = 360\text{ns}$, while
- ▶ the throughput is $1/120\text{ns} = 8.33 \times 10^6 \text{ operations/s}$.

Notes:

Sequential use of combinatorial stages \approx pipelining (9)

Example: a pipelined (logarithmic) shift

Circuit

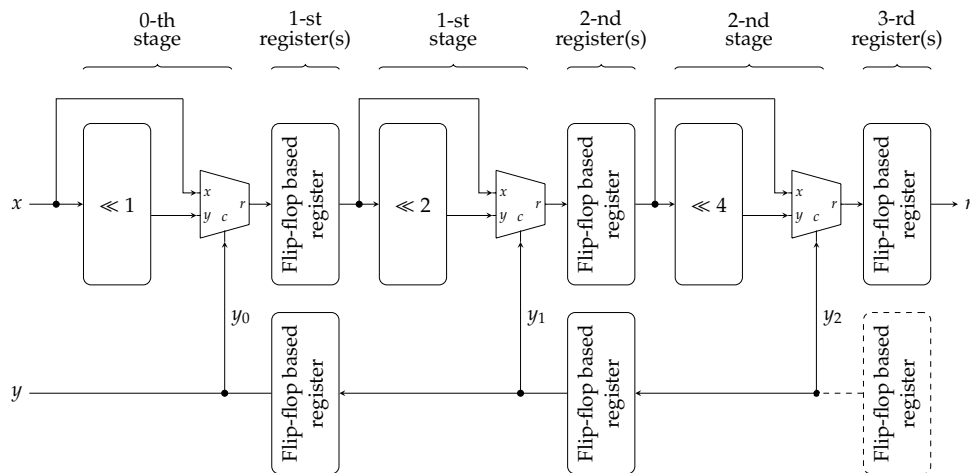


Notes:

Sequential use of combinatorial stages \approx pipelining (9)

Example: a pipelined (logarithmic) shift

Circuit



Notes:

Conclusions

► Take away points:

1. In a sense, pipelining is a classic trade-off, i.e.,

- ve: : higher area
- ve: : higher latency
- +ve: : higher throughput

2. Modulo some caveats, e.g.,

- need to balance stages, and
- need to keep the pipeline full,

pipelining is a *general-purpose* technique: it can be applied, for example, to execution of *instructions* by a micro-processor ...

Notes:

Conclusions

► Example: instruction pipelining [1].

- instruction execution follows a **fetch-decode-execute** cycle,
- but *various* approaches can implement this behaviour, e.g.,

Fetch, Decode, Execute, Commit

Notes:

1. a non-pipelined design.

Conclusions

► Example: instruction pipelining [1].

- instruction execution follows a **fetch-decode-execute** cycle,
- but *various* approaches can implement this behaviour, e.g.,




1. a non-pipelined design,
2. an **in-order** pipelined design.

Notes:

Conclusions

► Example: instruction pipelining [1].

- instruction execution follows a **fetch-decode-execute** cycle,
- but *various* approaches can implement this behaviour, e.g.,




1. a non-pipelined design,
2. an **in-order** pipelined design,
3. an **out-of-order** pipelined design.

Notes:

Conclusions

► Example: instruction pipelining [1].

- instruction execution follows a **fetch-decode-execute** cycle,
- but *various* approaches can implement this behaviour, e.g.,




1. a non-pipelined design,
2. an **in-order** pipelined design,
3. an **out-of-order** pipelined design,
4. a **Simultaneous MultiThreaded (SMT)** [3] pipelined design.

Notes:

Conclusions

► Example: instruction pipelining [1].

- instruction execution follows a **fetch-decode-execute** cycle,
- but *various* approaches can implement this behaviour, e.g.,



1. a non-pipelined design,
 2. an **in-order** pipelined design,
 3. an **out-of-order** pipelined design,
 4. a **Simultaneous MultiThreaded (SMT)** [3] pipelined design,
- or combinations thereof (e.g., out-of-order *and* SMT).

Notes:

Additional Reading

- ▶ *Wikipedia: Throughput*. URL: <http://en.wikipedia.org/wiki/Throughput>.
- ▶ *Wikipedia: Pipeline*. URL: [http://en.wikipedia.org/wiki/Pipeline_\(computing\)](http://en.wikipedia.org/wiki/Pipeline_(computing)).

Notes:

References

- [1] *Wikipedia: Instruction pipeline*. URL: http://en.wikipedia.org/wiki/Instruction_pipeline (see pp. 63, 65, 67, 69, 71).
- [2] *Wikipedia: Pipeline*. URL: [http://en.wikipedia.org/wiki/Pipeline_\(computing\)](http://en.wikipedia.org/wiki/Pipeline_(computing)) (see p. 73).
- [3] *Wikipedia: Simultaneous multithreading*. URL: http://en.wikipedia.org/wiki/Simultaneous_multithreading (see pp. 63, 65, 67, 69, 71).
- [4] *Wikipedia: Throughput*. URL: <http://en.wikipedia.org/wiki/Throughput> (see p. 73).
- [1] *Wikipedia: Instruction pipeline*. URL: http://en.wikipedia.org/wiki/Instruction_pipeline (see pp. 63, 65, 67, 69, 71).
- [2] *Wikipedia: Pipeline*. URL: [http://en.wikipedia.org/wiki/Pipeline_\(computing\)](http://en.wikipedia.org/wiki/Pipeline_(computing)) (see p. 73).
- [3] *Wikipedia: Simultaneous multithreading*. URL: http://en.wikipedia.org/wiki/Simultaneous_multithreading (see pp. 63, 65, 67, 69, 71).
- [4] *Wikipedia: Throughput*. URL: <http://en.wikipedia.org/wiki/Throughput> (see p. 73).

Notes: