

Intro. to Computer Architecture

Daniel Page

Department of Computer Science,
University Of Bristol,
Merchant Venturers Building,
Woodland Road,
Bristol, BS8 1UB. UK.
(csdsp@bristol.ac.uk)

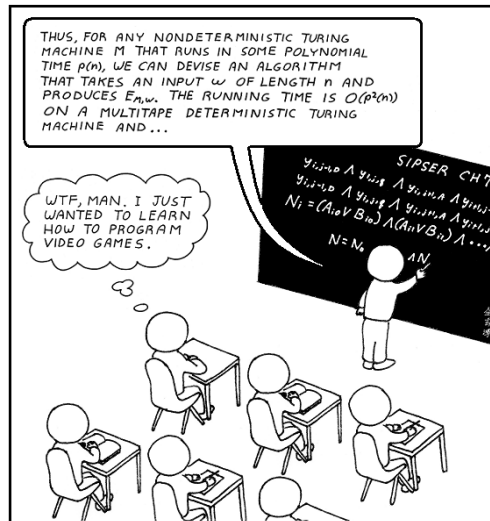
January 9, 2018

Keep in mind there are *two* PDFs available (of which this is the latter):

1. a PDF of examinable material used as lecture slides, and
2. a PDF of non-examinable, extra material:
 - ▶ the associated notes page may be pre-populated with extra, written explanation of material covered in lecture(s), plus
 - ▶ anything with a “grey’ed out” header/footer represents extra material which is useful and/or interesting but out of scope (and hence not covered).

Notes:

Notes:



<http://abstrusegoose.com/206>

© Daniel Page (dp@cs.bristol.ac.uk)
Intro. to Computer Architecture

git # 3b6f641 @ 2018-01-09



► (Fairly) reasonable **question(s)**:

1. "I thought this was CS, not Maths!", and
2. "why does this unit duplicate material in *other* units?".

Notes:

Notes:

► (Fairly) reasonable **question(s)**:

1. “I thought this was CS, not Maths!” and
2. “why does this unit duplicate material in *other* units?”.

► **Answer**: it isn’t, and it doesn’t (well, not *too* much) ... note

- theoretical concepts, e.g.,

$$\text{axiomatic manipulation} \leadsto \begin{cases} \text{optimisation} \\ \text{universality} \end{cases}$$

often have significant practical motivations or implications,

- it’s perfectly reasonable to utilise **Electronic Design Automation (EDA)** [3] tools, and
- Boolean algebra has wider application than hardware design.

Notes:

Boolean algebra: theory ≠ practice (1)

Axiomatic manipulation \leadsto optimisation

► **Question**: simplify the Boolean expression

$$(\neg(a \vee b) \wedge \neg(c \vee d \vee e)) \vee \neg(a \vee b)$$

into a form that contains the fewest operators possible.

Notes:

- **Question:** simplify the Boolean expression

$$(\neg(a \vee b) \wedge \neg(c \vee d \vee e)) \vee \neg(a \vee b)$$

into a form that contains the fewest operators possible.

- **Solution #1:** less steps.

$$\begin{aligned} & (\neg(a \vee b) \wedge \neg(c \vee d \vee e)) \vee \neg(a \vee b) \\ = & \neg(a \vee b) \vee (\neg(a \vee b) \wedge \neg(c \vee d \vee e)) \quad (\text{commutativity}) \\ = & \neg(a \vee b) \quad (\text{absorption}) \end{aligned}$$

Notes:

- **Question:** simplify the Boolean expression

$$(\neg(a \vee b) \wedge \neg(c \vee d \vee e)) \vee \neg(a \vee b)$$

into a form that contains the fewest operators possible.

- **Solution #2:** more steps.

$$\begin{aligned} & (\neg(a \vee b) \wedge \neg(c \vee d \vee e)) \vee \neg(a \vee b) \\ = & ((\neg a \wedge \neg b) \wedge \neg(c \vee d \vee e)) \vee \neg(a \vee b) \quad (\text{de Morgan}) \\ = & ((\neg a \wedge \neg b) \wedge (\neg c \wedge \neg d \wedge \neg e)) \vee \neg(a \vee b) \quad (\text{de Morgan}) \\ = & ((\neg a \wedge \neg b) \wedge (\neg c \wedge \neg d \wedge \neg e)) \vee (\neg a \wedge \neg b) \quad (\text{de Morgan}) \\ = & (\neg a \wedge \neg b) \vee ((\neg a \wedge \neg b) \wedge (\neg c \wedge \neg d \wedge \neg e)) \quad (\text{commutativity}) \\ = & (\neg a \wedge \neg b) \quad (\text{absorption}) \\ = & \neg(a \vee b) \quad (\text{de Morgan}) \end{aligned}$$

Notes:

- **Question:** simplify the Boolean expression

$$(a \wedge b \wedge c) \vee (\neg a \wedge b) \vee (a \wedge b \wedge \neg c)$$

into a form that contains the fewest operators possible.

Notes:

- **Question:** simplify the Boolean expression

$$(a \wedge b \wedge c) \vee (\neg a \wedge b) \vee (a \wedge b \wedge \neg c)$$

into a form that contains the fewest operators possible.

- **Solution:**

$$\begin{aligned} & (a \wedge b \wedge c) \vee (\neg a \wedge b) \vee (a \wedge b \wedge \neg c) \\ = & (a \wedge b \wedge c) \vee (a \wedge b \wedge \neg c) \vee (\neg a \wedge b) && \text{(commutativity)} \\ = & (a \wedge b) \wedge (c \vee \neg c) \vee (\neg a \wedge b) && \text{(distribution)} \\ = & (a \wedge b) \wedge 1 \vee (\neg a \wedge b) && \text{(inverse)} \\ = & (a \wedge b) \vee (\neg a \wedge b) && \text{(identity)} \\ = & b \wedge (a \vee \neg a) && \text{(distribution)} \\ = & b \wedge 1 && \text{(inverse)} \\ = & b && \text{(identity)} \end{aligned}$$

Notes:

Quote

If I designed a computer with 200 chips, I tried to design it with 150. And then I would try to design it with 100. I just tried to find every trick I could in life to design things real tiny.

– Wozniak

Quote

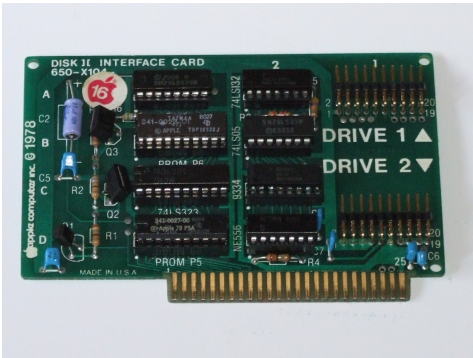
So I took 20 chips off their board; I bypassed 20 of their chips.

– Wozniak

Notes:

- The quotes relate to design and implementation of a (floppy) disk controller for the Apple II computer (circa 1977); there is an obvious focus on efficiency, which is credited as allowing the controller to be commercially viable. A detailed overview of the overarching anecdote is available via https://en.wikipedia.org/wiki/Disk_II
- or
- <http://apple2history.org/history/ah05/>
- The moral is that, in reality, “it works”, while important, may not be good enough: meeting various other (market-driven) quality metrics (e.g., efficiency, physical size, power consumption, etc.) is often vital rather than simply attractive.

Notes:



- ▶ We *can* add to our existing suite of operators (the result is often termed a **derived operator**), e.g.,
 - ▶ “NOT-AND” or **NAND**, st.

$$x \overline{\wedge} y \equiv \neg(x \wedge y)$$

so

x	y	$x \overline{\wedge} y$
0	0	1
0	1	1
1	0	1
1	1	0

and

- ▶ “NOT-OR” or **NOR**, st.

$$x \overline{\vee} y \equiv \neg(x \vee y)$$

so

x	y	$x \overline{\vee} y$
0	0	1
0	1	0
1	0	0
1	1	0

Notes:

Boolean algebra: theory ≠ practice (6)

Axiomatic manipulation → universality

- ▶ **Question:** haven’t we already got enough ... this is *already* quite difficult!
- ▶ **Answer:** NAND and NOR turn out to be **functionally complete** (or **universal**), e.g.,

$$\begin{aligned} \neg x &\equiv x \overline{\wedge} x \\ x \wedge y &\equiv (x \overline{\wedge} y) \overline{\wedge} (x \overline{\wedge} y) \\ x \vee y &\equiv \neg x \overline{\wedge} \neg y \equiv (x \overline{\wedge} x) \overline{\wedge} (y \overline{\wedge} y) \end{aligned}$$

which we can prove via

x	y	$x \overline{\wedge} y$	$x \overline{\wedge} x$	$y \overline{\wedge} y$	$(x \overline{\wedge} y) \overline{\wedge} (x \overline{\wedge} y)$	$(x \overline{\wedge} x) \overline{\wedge} (y \overline{\wedge} y)$
0	0	1	1	1	0	0
0	1	1	1	0	0	1
1	0	1	0	1	0	1
1	1	0	0	0	1	1

Notes:

- **Question:** haven't we already got enough ... this is *already* quite difficult!
- **Answer:** NAND and NOR turn out to be **functionally complete** (or **universal**), e.g.,

$$\begin{aligned} \neg x &\equiv x \bar{\wedge} x \\ x \wedge y &\equiv (x \bar{\wedge} y) \bar{\wedge} (x \bar{\wedge} y) \\ x \vee y &\equiv \neg x \bar{\wedge} \neg y \equiv (x \bar{\wedge} x) \bar{\wedge} (y \bar{\wedge} y) \end{aligned}$$

which we can prove via

x	y	$x \bar{\wedge} y$	$x \bar{\wedge} x$	$y \bar{\wedge} y$	$(x \bar{\wedge} y) \bar{\wedge} (x \bar{\wedge} y)$	$(x \bar{\wedge} x) \bar{\wedge} (y \bar{\wedge} y)$
0	0	1	1	1	0	0
0	1	1	1	0	0	1
1	0	1	0	1	0	1
1	1	0	0	0	1	1

- **Eureka:** computation of *any* Boolean function can be expressed using *one* simple building block component.

Notes:

- **Question:** translate

$$x \wedge (y \vee z)$$

into a version using NAND only.

Notes:

- **Question:** translate

$$x \wedge (y \vee z)$$

into a version using NAND only.

- **Solution #1:** apply the identities *naively* to get

$$\begin{aligned} & x \wedge (y \vee z) \\ = & x \wedge ((y \overline{\wedge} y) \overline{\wedge} (z \overline{\wedge} z)) \\ = & (x \overline{\wedge} ((y \overline{\wedge} y) \overline{\wedge} (z \overline{\wedge} z))) \overline{\wedge} (x \overline{\wedge} ((y \overline{\wedge} y) \overline{\wedge} (z \overline{\wedge} z))) \end{aligned}$$

Notes:

- **Question:** translate

$$x \wedge (y \vee z)$$

into a version using NAND only.

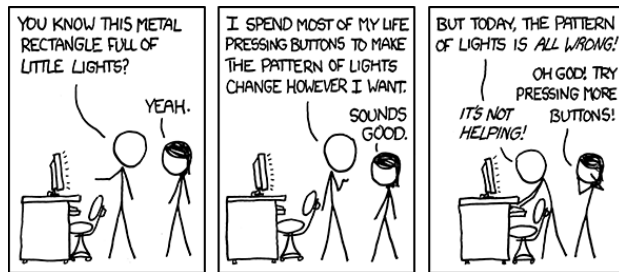
- **Solution #2:** apply the identities *intelligently* to get

$$\begin{aligned} & x \wedge (y \vee z) \\ = & x \wedge ((y \overline{\wedge} y) \overline{\wedge} (z \overline{\wedge} z)) \\ = & t \overline{\wedge} t \end{aligned}$$

where $t = x \overline{\wedge} ((y \overline{\wedge} y) \overline{\wedge} (z \overline{\wedge} z))$ is a common sub-expression [2].

Notes:

- ... which leads neatly to lab. worksheet #1:



<http://xkcd.com/722/>

Boolean algebra: theory \neq practice (9)

Utilising design automation tools

Listing

```
1 from sympy import *
2 from sympy.logic.boolalg import simplify_logic
3
4 a, b, c, d, e = symbols("a b c d e")
5
6 f = (~ (a | b) & ~(c | d | e)) | ~(a | b)
7
8 print f
9 print simplify_logic(f)
10 print
11
12 f = (a & b & c) | (~a & b) | (a & b & ~c)
13
14 print f
15 print simplify_logic(f)
16 print
```

Notes:

- Note that this example can be explored online via <http://live.sympy.org>
- The underlying point here is that it is often fine to use a computer to help, but only if you really understand the underlying theory: for more complex designs, this becomes a more important of course.
- Seymour Cray (a *supercomputer* architect) produced some interesting anecdotes about use of CAD tools in his design process: see for example

http://en.wikipedia.org/wiki/Seymour_Cray

Definition

Boolean algebra plays a role in many programming constructs:

$r \text{ is } x$	\equiv	$r = x$	\cong	$r == x$	\cong	$r = x$
$r \text{ is NOT } x$	\equiv	$r = \neg x$	\cong	$r != x$	\cong	$r = \sim x$
$r \text{ is } x \text{ NAND } y$	\equiv	$r = x \overline{\wedge} y$	\cong	$r != x \ \&\& \ y$	\cong	$r = \sim(x \ \& \ y)$
$r \text{ is } x \text{ NOR } y$	\equiv	$r = x \overline{\vee} y$	\cong	$r != x \ \ y$	\cong	$r = \sim(x \ \ y)$
$r \text{ is } x \text{ AND } y$	\equiv	$r = x \wedge y$	\cong	$r == x \ \&\& \ y$	\cong	$r = x \ \& \ y$
$r \text{ is } x \text{ OR } y$	\equiv	$r = x \vee y$	\cong	$r == x \ \ y$	\cong	$r = x \ \ y$
$r \text{ is } x \text{ XOR } y$	\equiv	$r = x \oplus y$			\cong	$r = x \ ^ \ y$

Note that the latter columns capture

- ▶ **decisional** operators, e.g., $\&\& : \mathbb{B} \times \mathbb{B} \rightarrow \mathbb{B}$, and
- ▶ **computational** operators, e.g., $\& : \mathbb{B}^n \times \mathbb{B}^n \rightarrow \mathbb{B}^n$.

Notes:

Listing

```
1 void condition_v1( int a, int b, int c, int d ) {
2   if( ( ( a + b ) > c ) && ( ( d % 2 ) == 0 ) ) {
3     //      ( a + b is   > c ) and    ( d is   even )
4
5     printf( "if -> true   : %d %d %d %d\n", a, b, c, d );
6   }
7   else {
8     //      not ( ( a + b is   > c ) and    ( d is   even ) )
9     //
10    // = not   ( a + b is   > c ) or  not ( d is   even )
11    // =      ( a + b isn't > c ) or  ( d isn't even )
12    // =      ( a + b is   <= c ) or  ( d is   odd  )
13
14    printf( "if -> false  : %d %d %d %d\n", a, b, c, d );
15  }
16 }
```

Notes:

Listing

```
1 void condition_v2( int a, int b, int c, int d ) {  
2   if( ( ( a + b ) > c ) && ( d = 3 ) ) {  
3     printf( "if -> true   : %d %d %d %d\n", a, b, c, d );  
4   }  
5   else {  
6     printf( "if -> false : %d %d %d %d\n", a, b, c, d );  
7   }  
8 }
```

Notes:

- **Question:** imagine we have a Boolean expression (or circuit)

$$f(x_0, x_1, \dots x_{n-1})$$

and manipulate it into a different form

$$g(x_0, x_1, \dots x_{n-1})$$

e.g., optimise it: we want to know if there is a bug in g .

Notes:

- **Question:** imagine we have a Boolean expression (or circuit)

$$f(x_0, x_1, \dots x_{n-1})$$

and manipulate it into a different form

$$g(x_0, x_1, \dots x_{n-1})$$

e.g., optimise it: we want to know if there is a bug in g .

- **Solution(s):**

1. *prove* that f and g are equivalent,
2. try all 2^n possible assignments, and see if f ever differs from g , or
3. use a (carefully defined) instance of **SAT**.

Notes:

Definition

The **Boolean satisfiability problem** (or **SAT**) is a decision problem. Consider some Boolean function

$$f(x_0, x_1, \dots x_{n-1}).$$

which defines a SAT **instance**. The problem is to decide whether or not an assignment to said variables (i.e., $x_i \in \{0, 1\}$ for $0 \leq i < n$) exists st. f is **satisfiable** (i.e., we have $f(x_0, x_1, \dots x_{n-1}) = 1$).

Notes:

- SAT is NP-complete [7, 8] meaning that
 1. no known algorithm will efficiently solve every SAT instance, but
 2. other NP problems can be expressed as SAT instances.

► **Question:** decide whether

$$f(a, b) = (b \wedge \neg a) \vee (a \wedge \neg b)$$

\equiv

$$g(a, b) = (a \vee b) \wedge \neg(a \wedge b).$$

Notes:

- The idea of the SAT instance is to use how XOR works: what we're saying is that if h is satisfiable, then we have an assignment to the variables st. the result of f and g differ (XOR produces 1 as output if one (and only one) of the inputs are 1, and 0 otherwise).
- One underlying idea here is that each approach is valid, but also might be advantageous or disadvantageous when used in a particular context: the use of axiomatic manipulation does not *tell* us anything (e.g., why f and g differ, if they do), for example, and brute-force enumeration will clearly become intractable very quickly (as n grows). The other is that greater understanding of the relationship between theory and practice allows better results in both: here, in a rough sense,
 - in theory, SAT helps us understand *what* we can compute, while
 - in practice, SAT helps us understand *how* we compute it.

► **Question:** decide whether

$$f(a, b) = (b \wedge \neg a) \vee (a \wedge \neg b)$$

\equiv

$$g(a, b) = (a \vee b) \wedge \neg(a \wedge b).$$

► **Solution #1:** manipulation via axioms, e.g.,

$$\begin{aligned}
 f(a, b) &= (b \wedge \neg a) && \vee && (a \wedge \neg b) \\
 &= (b \wedge \neg a) \vee 0 && \vee && (a \wedge \neg b) \vee 0 && \text{(identity)} \\
 &= (b \wedge \neg a) \vee (b \wedge \neg b) && \vee && (a \wedge \neg b) \vee (a \wedge \neg a) && \text{(inverse)} \\
 &= (b \wedge (\neg a \vee \neg b)) && \vee && (a \wedge (\neg b \vee \neg a)) && \text{(distribution)} \\
 &= ((\neg a \vee \neg b) \wedge b) && \vee && ((\neg a \vee \neg b) \wedge a) && \text{(commutativity)} \\
 &= (\neg a \vee \neg b) \wedge (a \vee b) && && && \text{(distribution)} \\
 &= (a \vee b) \wedge (\neg a \vee \neg b) && && && \text{(commutativity)} \\
 &= (a \vee b) \wedge \neg(a \wedge b) && && && \text{(de Morgan)} \\
 &= g(a, b)
 \end{aligned}$$

Notes:

- The idea of the SAT instance is to use how XOR works: what we're saying is that if h is satisfiable, then we have an assignment to the variables st. the result of f and g differ (XOR produces 1 as output if one (and only one) of the inputs are 1, and 0 otherwise).
- One underlying idea here is that each approach is valid, but also might be advantageous or disadvantageous when used in a particular context: the use of axiomatic manipulation does not *tell* us anything (e.g., why f and g differ, if they do), for example, and brute-force enumeration will clearly become intractable very quickly (as n grows). The other is that greater understanding of the relationship between theory and practice allows better results in both: here, in a rough sense,
 - in theory, SAT helps us understand *what* we can compute, while
 - in practice, SAT helps us understand *how* we compute it.

► **Question:** decide whether

$$f(a, b) = (b \wedge \neg a) \vee (a \wedge \neg b)$$

\equiv

$$g(a, b) = (a \vee b) \wedge \neg(a \wedge b).$$

► **Solution #2:** brute-force enumeration, i.e.,

a	b	$b \wedge \neg a$	$a \wedge \neg b$	$f(a, b)$	$a \vee b$	$\neg(a \wedge b)$	$g(a, b)$
0	0	0	0	0	0	1	0
0	1	1	0	1	1	1	1
1	0	0	1	1	1	1	1
1	1	0	0	0	1	0	0

Notes:

- The idea of the SAT instance is to use how XOR works: what we're saying is that if h is satisfiable, then we have an assignment to the variables st. the result of f and g *differ* (XOR produces 1 as output if one (and only one) of the inputs are 1, and 0 otherwise).
- One underlying idea here is that each approach is valid, but also might be advantageous or disadvantageous when used in a particular context: the use of axiomatic manipulation does not *tell* us anything (e.g., why f and g differ, if they do), for example, and brute-force enumeration will clearly become intractable very quickly (as n grows). The other is that greater understanding of the relationship between theory and practice allows better results in both: here, in a rough sense,
 - in theory, SAT helps us understand *what* we can compute, while
 - in practice, SAT helps us understand *how* we compute it.

► **Question:** decide whether

$$f(a, b) = (b \wedge \neg a) \vee (a \wedge \neg b)$$

\equiv

$$g(a, b) = (a \vee b) \wedge \neg(a \wedge b).$$

► **Solution #3:** define

$$h(x_0, x_1, \dots, x_{n-1}) = f(x_0, x_1, \dots, x_{n-1}) \oplus g(x_0, x_1, \dots, x_{n-1})$$

and use this as a SAT instance: if the SAT instance is satisfiable, this yields a **test vector** we can use to debug g .

Notes:

- The idea of the SAT instance is to use how XOR works: what we're saying is that if h is satisfiable, then we have an assignment to the variables st. the result of f and g *differ* (XOR produces 1 as output if one (and only one) of the inputs are 1, and 0 otherwise).
- One underlying idea here is that each approach is valid, but also might be advantageous or disadvantageous when used in a particular context: the use of axiomatic manipulation does not *tell* us anything (e.g., why f and g differ, if they do), for example, and brute-force enumeration will clearly become intractable very quickly (as n grows). The other is that greater understanding of the relationship between theory and practice allows better results in both: here, in a rough sense,
 - in theory, SAT helps us understand *what* we can compute, while
 - in practice, SAT helps us understand *how* we compute it.

Listing

```
1 from sympy import *
2 from sympy.logic.inference import satisfiable
3
4 a, b, c, d, e = symbols( "a b c d e" )
5
6 f = ( b & ~a ) | ( a & ~b )
7 g = ( a | b ) & ~( a & b )
8
9 print f
10 print g
11 print satisfiable( f ^ g )
12 print
```

Notes:

- Note that this example can be explored online via <http://live.sympy.org>

Conclusions

Quote

In theory, there is no difference between theory and practice. But, in practice, there is.

– van de Snepscheut (http://en.wikiquote.org/wiki/Jan_L._A._van_de_Snepscheut)

► **Take away points:** being pragmatic, it should be clear

1. “it works” \neq “it works *well*”,
2. using automation is fine *iff.* you know the underlying theory,
3. using brute-force is fine *iff.* you know the underlying theory,
4. Boolean algebra > Boolean axioms: concepts that *seem* of interest in theory alone, *can* be important if/when applied in practice,
5. computer architecture > hardware: Boolean algebra can also explain/support concepts in software.

Notes:

Additional Reading

- ▶ *Wikipedia: Boolean algebra*. URL: http://en.wikipedia.org/wiki/Boolean_algebra.
- ▶ D. Page. “Chapter 1: Mathematical preliminaries”. In: *A Practical Introduction to Computer Architecture*. 1st ed. Springer-Verlag, 2009.
- ▶ W. Stallings. “Chapter 11: Digital logic”. In: *Computer Organisation and Architecture*. 9th ed. Prentice-Hall, 2013.
- ▶ A.S. Tanenbaum and T. Austin. “Section 3.1: Gates and Boolean algebra”. In: *Structured Computer Organisation*. 6th ed. Prentice-Hall, 2012.

Notes:

References

- [1] *Wikipedia: Boolean algebra*. URL: http://en.wikipedia.org/wiki/Boolean_algebra (see p. 65).
- [2] *Wikipedia: Common sub-expression elimination*. URL: http://en.wikipedia.org/wiki/Common_subexpression_elimination (see pp. 33, 35).
- [3] *Wikipedia: Electronic Design Automation (EDA)*. URL: http://en.wikipedia.org/wiki/Electronic_design_automation (see pp. 7, 9).
- [4] D. Page. “Chapter 1: Mathematical preliminaries”. In: *A Practical Introduction to Computer Architecture*. 1st ed. Springer-Verlag, 2009 (see p. 65).
- [5] W. Stallings. “Chapter 11: Digital logic”. In: *Computer Organisation and Architecture*. 9th ed. Prentice-Hall, 2013 (see p. 65).
- [6] A.S. Tanenbaum and T. Austin. “Section 3.1: Gates and Boolean algebra”. In: *Structured Computer Organisation*. 6th ed. Prentice-Hall, 2012 (see p. 65).
- [7] S. Cook. “The complexity of theorem proving procedures”. In: *ACM Symposium on Theory of Computing (STOC)*. 1971, pp. 151–158 (see p. 52).
- [8] L. Levin. “Universal search problems”. In: *Problems of Information Transmission* 9.3 (1973), pp. 265–266 (see p. 52).

Notes: