

# Intro. to Computer Architecture

Daniel Page

Department of Computer Science,  
University Of Bristol,  
Merchant Venturers Building,  
Woodland Road,  
Bristol, BS8 1UB. UK.  
([csdsp@bristol.ac.uk](mailto:csdsp@bristol.ac.uk))

January 9, 2018

Keep in mind there are *two* PDFs available (of which this is the latter):

1. a PDF of examinable material used as lecture slides, and
2. a PDF of non-examinable, extra material:
  - ▶ the associated notes page may be pre-populated with extra, written explanation of material covered in lecture(s), plus
  - ▶ anything with a “grey’ed out” header/footer represents extra material which is useful and/or interesting but out of scope (and hence not covered).

Notes:

Notes:

► An  $n$ -bit register based on latches (resp. flip-flops) has two **limitations**:

1. each latch (resp. flip-flop) in the register needs a relatively large number of transistors, which limits the viable capacity (i.e.,  $n$ ), and
2. the register is not addressable, i.e.,
  - an **address** (or **index**) allows dynamic rather than static reference to some stored datum, so
  - by rough analogy to a C program


Listing	
1	<code>int A0, A1, A2, A3;</code>
2	
3	<code>A0 = 0;</code>
4	<code>A1 = 0;</code>
5	<code>A2 = 0;</code>
6	<code>A3 = 0;</code>

Listing	
1	<code>int A[ 4 ];</code>
2	
3	<code>A[ 0 ] = 0;</code>
4	<code>A[ 1 ] = 0;</code>
5	<code>A[ 2 ] = 0;</code>
6	<code>A[ 3 ] = 0;</code>

we *have* the left-hand side, but we *want* the right-hand side.

Notes:

► **Solution**: a **memory** component, i.e.,



st.

- MEM has a capacity of  $n = 2^{n'}$  addressable words,
- each such word is  $w$  bits (where  $n \gg w$ ).

Notes:

- ▶ There are various ways to classify a given memory component, e.g.,
  1. volatility:
    - ▶ **volatile**, meaning the content is lost when the component is powered-off, or
    - ▶ **non-volatile**, meaning the content is retained even after the component is powered-off.
  2. interface type:
    - ▶ **synchronous**, where a clock or pre-determined timing information synchronises steps, or
    - ▶ **asynchronous**, where a protocol synchronises steps.
  3. access type:
    - ▶ random versus constrained (e.g., sequential) access to content,
    - ▶ **Random Access Memory (RAM)** which we can read from *and* write to, and
    - ▶ **Read Only Memory (ROM)** which, as suggested by the name, supports reads only.
  4. ...

but we'll focus on a volatile, synchronous RAM.


Notes:

## An Aside: History



- ▶ The EDSAC used **delay line** memory, where the rough idea is:
  - ▶ Each “line” is a tube of mercury (or something else in which sound waves propagate fairly slowly).
  - ▶ Put a speaker at one end to store sound waves into the line, and a microphone at the other to read them out.
  - ▶ Values are stored in the sense the corresponding waves take time to propagate; when they get to one end they are either replaced or fed back into the other.
- ▶ This is **sequential access** (cf. **random access**): you need to *wait* for the data you want to appear!

Notes:



- ▶ The Whirlwind used **magnetic-core** memory, where the rough idea is:
  - ▶ The memory is a matrix of small magnetic rings, or “cores”, which can be magnetically polarised to store values.
  - ▶ Wires are threaded through the cores to control them, i.e., to store or read values.
  - ▶ The magnetic polarisation is retained, so core memory is non-volatile!
- ▶ You might still hear main memory termed **core memory** (cf. **core dump**) which is a throw-back to this technology.

[http://en.wikipedia.org/wiki/File:Project\\_Whirlwind\\_-\\_core\\_memory,\\_circa\\_1951\\_-\\_detail\\_1.JPG](http://en.wikipedia.org/wiki/File:Project_Whirlwind_-_core_memory,_circa_1951_-_detail_1.JPG)

## Low-level Implementation (1)

Definition	Definition
<b>Static RAM (SRAM)</b> is <ul style="list-style-type: none"><li>▶ manufacturable in lower densities (i.e., smaller capacity),</li><li>▶ more expensive to manufacture,</li><li>▶ fast(er) access time (resp. lower access latency),</li><li>▶ easy(er) to interface with,</li><li>▶ ideal for latency-optimised contexts, e.g., as <b>cache memory</b>.</li></ul>	<b>Dynamic RAM (DRAM)</b> is <ul style="list-style-type: none"><li>▶ manufacturable in higher densities (i.e., larger capacity),</li><li>▶ less expensive to manufacture,</li><li>▶ slow(er) access time (resp. higher access latency),</li><li>▶ hard(er) to interface with,</li><li>▶ ideal for capacity-optimised contexts, e.g., as <b>main memory</b>.</li></ul>

Notes:


Notes:

## Low-level Implementation (2)

An SRAM cell

- Abstractly, an **SRAM cell** resembles two NOT gates:

### Circuit




Notes:

## Low-level Implementation (2)

An SRAM cell

- Abstractly, an **SRAM cell** resembles two NOT gates ...

### Circuit



Notes:

- ... concretely, we can fill in the NOT gates with the transistor-level equivalents; each “6T SRAM cell” requires 6 transistors (cf.  $\sim 20$  or so for a D-type latch).

- ▶ Note that:
  - ▶ The initial NOT-based circuit might look odd, but clearly has two stable states: either  $Q = 1$  and  $\neg Q = 0$ , or  $Q = 0$  and  $\neg Q = 1$ .
  - ▶ The transistors re-enforce each other; the state is maintained as long as the cell is powered-on.
  - ▶  $bl$  and  $\neg bl$  are termed the **bit lines**,  $wl$  is the **word line** which controls access to the state.
  - ▶ The pre-charging steps are managed by extra **bit line conditioning** logic which we gloss over from here on.
- ▶ So
  - ▶ to read the cell we pre-charge  $bl = 1$  and  $\neg bl = 1$  then set  $wl = 1$ , after which  $\neg bl$  (resp.  $bl$ ) is discharged if state is 1 (resp. 0), or
  - ▶ to write  $x$  into the cell we pre-charge  $bl = x$  and  $\neg bl = \neg x$  then set  $wl = 1$ , after which the state matches  $x$ .


Notes:

## Low-level Implementation (4)

A DRAM cell

- ▶ A **DRAM cell** is constructed using 1 transistor and a capacitor:

Circuit



Notes:

- ▶ Note that:
  - ▶ The state *decays* when the cell is read, *and* also over time (even if it's not read) due to leakage; this implies a need to **refresh** the cell periodically.
  - ▶ The capacitor holds a tiny charge: this must be amplified to use as a driver in whatever circuit uses the cell.
  - ▶ The speed at which the capacitor charges and discharges is (relatively) slow.
- ▶ So
  - ▶ to read the cell we set  $wl = 1$ , after which current flows (resp. does not flow) on  $bl$  if the capacitor is charged (resp. not charged) meaning the state is 1 (resp. 0), or
  - ▶ to write  $x$  into the cell we set  $bl = x$  then  $wl = 1$ , after which the capacitor charges (resp. discharges) and the state matches  $x$ .

Notes:

## High(er)-level Implementation (1)

Cells  $\rightsquigarrow$  Device

- ▶ A **memory device** is constructed from (roughly) three components
  1. a **memory array** (or matrix) of replicated cells with
    - ▶  $r$  rows, and
    - ▶  $c$  columnsmeaning a  $(r \cdot c)$ -cell capacity.
  2. a **row decoder** which given an address (de)activates associated cells in that row, and
  3. a **column decoder** which given an address (de)selects associated cells in that columnplus additional logic to allow use (depending on cell type), e.g.,
  1. **bit line conditioning** to ensure the bit lines are strong enough to be effective,
  2. **sense amplifiers** to ensure output from the array is usable.

Notes:

► A 1-bit **SRAM device** with  $n = 2^{n'}$  cells has a (somewhat) standard physical interface:

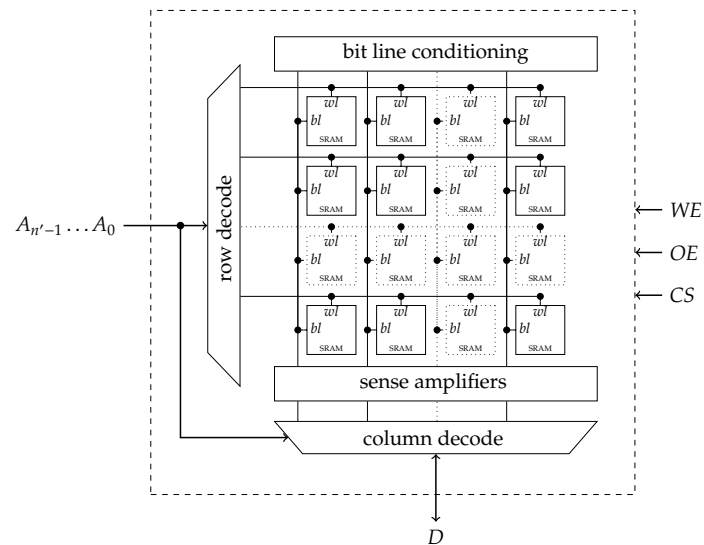
1. auxiliary pin(s) for power and so on,
2.  $D$ , a single 1-bit **data pin** (sometimes split into two separate  $D_{in}$  and  $D_{out}$  pins),
3.  $A$ , a collection of  $n'$  **address pins** where  $A_i$  is the  $i$ -th such pin,
4. a **Chip Select (CS)** pin, which enables the device,
5. a **Output Enable (OE)** pin, which signals the device is being read from, and
6. a **Write Enable (WE)** pin, which signals the device is being written to.

Notes:

## High(er)-level Implementation (3)

Cells  $\rightsquigarrow$  Device: An SRAM device

Example (an  $n$ -cell SRAM memory device)



Notes:



Algorithm (SRAM-READ)

Having performed the following steps

- drive the address onto  $A$ ,
- set  $WE = \text{false}$ ,  $OE = \text{true}$  and  $CS = \text{true}$

1-bit of data is read and made available on  $D$ , then we set  $CS = \text{false}$ .

Algorithm (SRAM-WRITE)

Having performed the following steps

- drive the address onto  $A$ ,
- drive the data onto  $D$ ,
- Set  $WE = \text{true}$ ,  $OE = \text{false}$  and  $CS = \text{true}$

1-bit of data is written, then we set  $CS = \text{false}$ .

Notes:

Notes:

ASI  
Austin Semiconductor, Inc.

SRAM  
MT5C1001  
Limited Availability

1M x 1 SRAM  
SRAM MEMORY ARRAY  
AVAILABLE AS MILITARY  
SPECIFICATIONS

- MIL-STD-883C
- MIL-STD-883

FEATURES

- High-Speed: 25, 25, 35, and 45
- Battery Backed: 2V data retention
- Low power standby
- Single +5V (±10%) Power Supply
- Easy memory expansion with CE1 and CE0 options
- All inputs and outputs are TTL compatible
- Three-state output

OPTIONS

MARKING

- Timing
- 25ns access -3D
- 25ns access -3S
- 35ns access -3S
- 45ns access -4S
- 55ns access -5S\*
- 70ns access -7D\*

- Packages
- Ceramic DIP (40 pins) C No. 109
- Ceramic LCC 48 No. 307
- Ceramic Flipchip F No. 303
- Ceramic QFP Q No. 301

Operating Temperature Ranges

- Industrial (-40°C to +85°C) IT
- Military (-55°C to +125°C) MT

- 2V data retention/low power L

GENERAL DESCRIPTION

The MT5C1001 employs low power, high performance silicon gate CMOS technology. Static design eliminates the need for external clocks or timing modes while CMOS circuitry reduces power consumption and provides for greater reliability.

For flexibility in high-speed memory applications, All other chip enable (CE) and output enable (OE) capability. These enhancements can place the outputs in High-Z for additional flexibility in system design. Writing to these devices is accomplished when write enable (WE) and CE inputs are both LOW. Reading is accomplished when WE assumes HIGH while CE1 and CE0 are LOW. The device offers a reduced power standby mode when disabled. This allows system designers to achieve low standby power requirements.

The "V" notation provides an approximate 50 percent reduction in CMOS standby current ( $I_{CCS}$ ) over the standard version.

All devices operate from a single +5V power supply and all inputs and outputs are fully TTL compatible.

For more products and information  
please visit our web site at  
[www.austinsemiconductor.com](http://www.austinsemiconductor.com)

1

ASI  
Austin Semiconductor, Inc.  
Rev. 1.07, 1/98

ASI  
Austin Semiconductor, Inc.

SRAM  
MT5C1001  
Limited Availability

FUNCTIONAL BLOCK DIAGRAM

TRUTHTABLE

MODE	CE1	WE	OUTPUT	POWER
STANDBY	H	X	High-Z	STANDBY
READ	L	H	D	ACTIVE
WRITE	L	L	DATA IN	ACTIVE

PIN ASSIGNMENTS

PIN	ASSIGNMENT
A0-A10	Address Inputs
CE1	Write Enable
CE0	Chip Enable
D	Data Input
D	Data Output
OE	Output Enable
WE	Write Enable
VCC	+5V Power Supply
VDD	Ground

2

ASI  
Austin Semiconductor, Inc.  
Rev. 1.07, 1/98

© Daniel Page ( [csdps@bristol.ac.uk](mailto:csdps@bristol.ac.uk) )  
Intro. to Computer Architecture

git # 3b6f641 @ 2018-01-09

University of  
BRISTOL

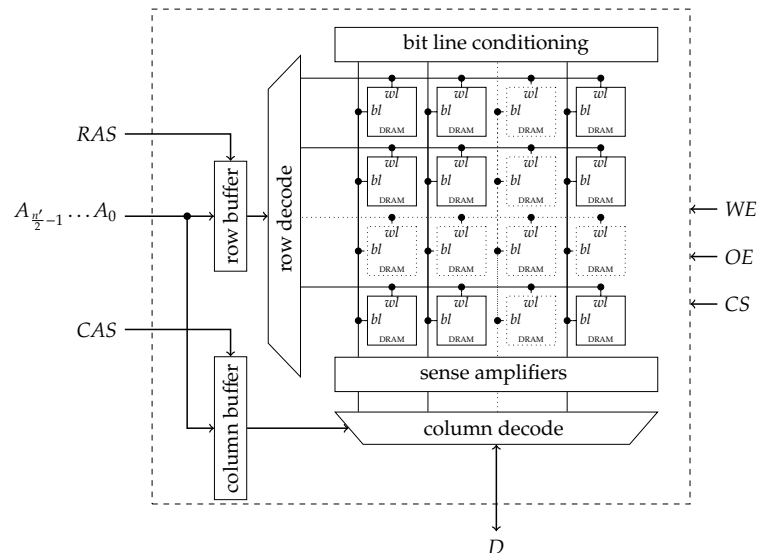
- ▶ A 1-bit **DRAM device** with  $n = 2^{n'}$  cells has a (somewhat) standard physical interface:
  1. auxiliary pin(s) for power and so on,
  2.  $D$ , a single 1-bit **data pin** (sometimes split into two separate  $D_{in}$  and  $D_{out}$  pins),
  3.  $A$ , a collection of  $\frac{n'}{2}$  **address pins** where  $A_i$  is the  $i$ -th such pin,
  4. a **Chip Select (CS)** pin, which enables the device,
  5. a **Output Enable (OE)** pin, which signals the device is being read from,
  6. a **Write Enable (WE)** pin, which signals the device is being written to,
  7. a **Row Address Strobe (RAS)**, which controls the row buffer, and
  8. a **Column Address Strobe (CAS)**, which controls the column buffer.
- ▶ Note there are typically *half* the number of address pins versus the same sized SRAM device:
  - ▶ the pins are multiplexed to form a full address,
  - ▶ since DRAM is more dense, this acts to manage the number of pins required.

Notes:

## High(er)-level Implementation (7)

Cells  $\rightsquigarrow$  Device: A DRAM device

### Example (a $n$ -cell DRAM memory device)



Notes:

Algorithm (DRAM-READ)	Algorithm (DRAM-WRITE)
<p>Having performed the following steps</p> <ol style="list-style-type: none"> <li>1. Drive the row address onto <i>addr</i>.</li> <li>2. Set <i>RAS</i> = <b>true</b>, which latches row address.</li> <li>3. Drive the column address onto <i>addr</i>.</li> <li>4. Set <i>CAS</i> = <b>true</b>, which latches column address.</li> <li>5. Set <i>WE</i> = <b>false</b>, <i>OE</i> = <b>true</b> and <i>CS</i> = <b>true</b>.</li> </ol> <p>1-bit of data is read and made available on <i>D</i>, and we set <i>CSRAS</i> = <i>CAS</i> = <b>false</b>.</p>	<p>Having performed the following steps</p> <ol style="list-style-type: none"> <li>1. Drive the row address onto <i>addr</i>.</li> <li>2. Set <i>RAS</i> = <b>true</b>, which latches row address.</li> <li>3. Drive the column address onto <i>addr</i>.</li> <li>4. Set <i>CAS</i> = <b>true</b>, which latches column address.</li> <li>5. Drive the data onto <i>data</i>.</li> <li>6. Set <i>WE</i> = <b>false</b>, <i>OE</i> = <b>true</b> and <i>CS</i> = <b>true</b>.</li> </ol> <p>1-bit of data is written, and we set <i>CSRAS</i> = <i>CAS</i> = <b>false</b>.</p>

- ▶ Plus we need logic to implement DRAM refresh:
  - ▶ To cope with decay of cell content, we periodically read *all* the cells.
  - ▶ This amounts to activating each row in turn; the latency of refresh can therefore be optimised by selecting smaller  $r$  (resp. larger  $c$ ) for the array.

[illegible][illegible]

Notes:

Notes:

- Externally, the configuration of a device is described as something like

$$\delta \times \omega \times \beta$$

(plus maybe some timing information) where

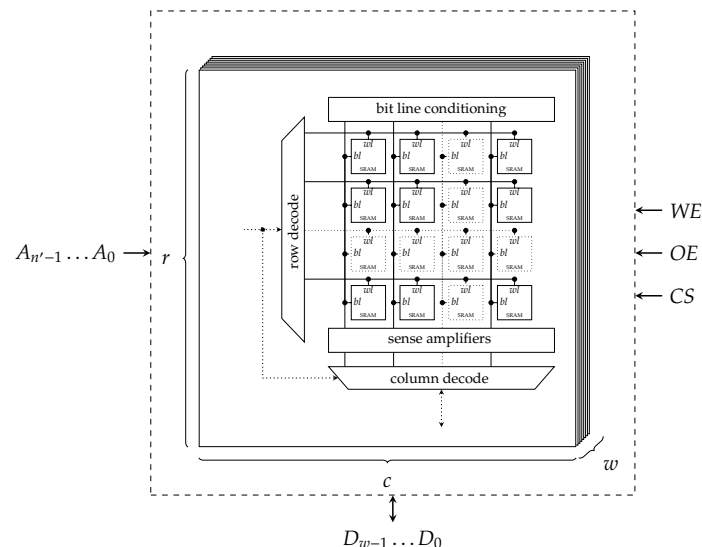
- $\delta$  relates to capacity, usually measured in (large multiples of) bits,
  - $\omega$  describes the width of words, measured in bits,
  - $\beta$  is the number of internal **logical banks**.
- Internally, this implies some organisational choices:
- for  $\omega > 1$ , we replicate the memory device internally to give  $\omega$  arrays (each copy relates to one bit of a  $\omega$ -bit word),
  - for the arrays,  $r$  and  $c$  can be selected to match physical requirements (e.g., to get “square” or “thin” arrays), and
  - for  $\beta > 1$ , each array is split into logical banks

where in the latter case, the goal is to distribute the range of addresses across multiple smaller banks.

Notes:

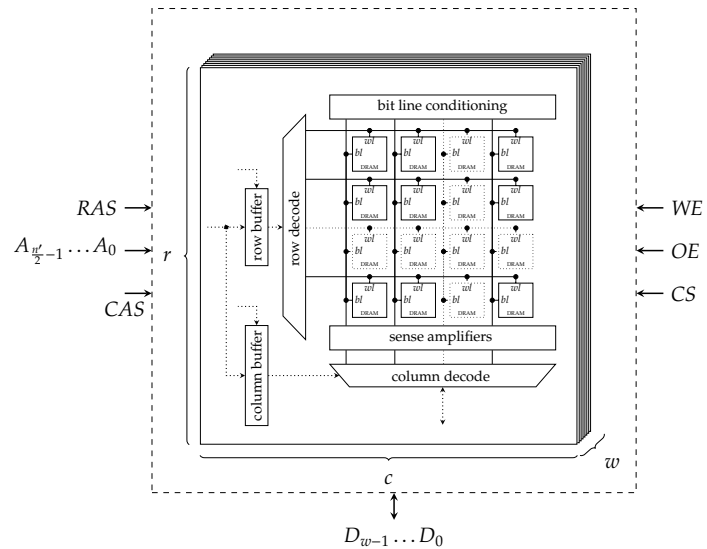
## High(er)-level Implementation (11) Cells $\rightsquigarrow$ Device

Example (from a 1-bit to  $w$ -bit SRAM device via replication)



Notes:

### Example (from a 1-bit to $w$ -bit DRAM device via replication)




Notes:

### High(er)-level Implementation (13) Cells $\rightsquigarrow$ Device

#### Example (a 512Mbit, 8-bit Device)

- ▶ A 64Mbit  $\times$  8  $\times$  1 internal configuration:
  - ▶ The total capacity is 64Mbit  $\cdot$  8  $\cdot$  1 = 512Mbit.
  - ▶ There is  $\beta = 1$  logical bank, consisting of  $\omega = 8$  arrays; each array has  $\delta = r \cdot c = 64 \cdot 1024 \cdot 1024$  cells.
  - ▶ Address  $x$  refers to cell  $x$  of each array in the bank; for example  $x = 42_{(10)}$  refers to cell 42 of each array, and therefore to an 8-bit word overall.



Notes:

### Example (a 512Mbit, 8-bit Device)

- ▶ A 32Mbit  $\times$  8  $\times$  2 internal configuration:


- ▶ The total capacity is 32Mbit  $\cdot$  8  $\cdot$  2 = 512Mbit.
- ▶ There are  $\beta = 2$  logical banks, each consisting of  $\omega = 8$  arrays; each array has  $\delta = r \cdot c = 32 \cdot 1024 \cdot 1024$  cells.
- ▶ Address  $x$  refers to

$$x \mapsto \begin{cases} \text{bank} & x \bmod \beta \\ \text{cell} & x \div \beta \end{cases}$$

of each array; for example,

1.  $x = 42_{(10)}$  refers to cell 21 of each array in bank 0, while
2.  $x = 43_{(10)}$  refers to cell 21 of each array in bank 1

each case therefore referring to 8-bit word overall.



Notes:

### Example (a 512Mbit, 8-bit Device)

- ▶ A 16Mbit  $\times$  8  $\times$  4 internal configuration:


- ▶ The total capacity is 16Mbit  $\cdot$  8  $\cdot$  4 = 512Mbit.
- ▶ There are  $\beta = 4$  logical banks, each consisting of  $\omega = 8$  arrays; each array has  $\delta = r \cdot c = 16 \cdot 1024 \cdot 1024$  cells.
- ▶ Address  $x$  refers to

$$x \mapsto \begin{cases} \text{bank} & x \bmod \beta \\ \text{cell} & x \div \beta \end{cases}$$

of each array; for example,


1.  $x = 42_{(10)}$  refers to cell 10 of each array in bank 2, while
2.  $x = 43_{(10)}$  refers to cell 21 of each array in bank 3

each case therefore referring to 8-bit word overall.



Notes:

### Example (an $n$ -cell ROM memory device)



Notes:

### High-level Implementation (1)

Device  $\leadsto$  Module

- ▶ To make devices easier to use (or integrate), they are typically packaged into a **memory module**:
    1. one or more memory devices, and
    2. an interface which controls access.
  - ▶ There are lots of package types; two dominate
    1. **Single Inline Memory Module (SIMM)**, which is (roughly) 1-sided, has less pins and a narrower word size, and
    2. **Dual Inline Memory Module (DIMM)**, which is (roughly) 2-sided, has more pins and a wider word size
- and are capable of housing different device types (e.g., EDO *or* SDRAM devices in a given DIMM package).

Notes:

- Externally, the configuration of a module is described as something like

$$\Delta \times \Omega$$

(plus maybe some timing information) where


- $\Delta$  relates to capacity, usually measured in (large multiplies of) bytes,
  - $\Omega$  describes the width of words, measured in bits.
- Internally, some organisational choices exist
    - usually  $\Omega > \omega$  so the module is “filled” using multiple devices to form one **physical bank**, and
    - depending on the module type, the devices are organised into one or more **ranks**

where in the former case, the goal is to distribute content relating to a single address across multiple smaller devices.

Notes:

### Example (a 4MB, 32-bit SIMM)

- A  $4\text{MB} \times 32$ , 1-device configuration:
  - The device has a capacity of  $1\text{Mbit} \cdot 32 \cdot 1 = 32\text{Mbit} = 4\text{MB}$  arranged into 32-bit words.
  - The bank, which has a 32-bit data-path, is filled by the single device: address  $x$  refers to the word  $\text{MEM}$ .



Notes:




## Example (a 4MB, 32-bit SIMM)

- ▶ A 4MB  $\times$  32, 2-device configuration:
  - ▶ Each device has a capacity of 1Mbit  $\cdot$  16  $\cdot$  1 = 16Mbit = 2MB arranged into 16-bit words.
  - ▶ There are 2 devices; the total capacity is 2  $\cdot$  2MB = 4MB.
  - ▶ The bank, which has a 32-bit data-path, is filled by merging the devices: address  $x$  refers to

$$x \mapsto \begin{cases} \text{bits } 15 \dots 0 & \text{of MEM from device 0} \\ \text{bits } 31 \dots 16 & \text{of MEM from device 1} \end{cases}$$

merging each 16-bit part into the word MEM.




Notes:

## Example (a 4MB, 32-bit SIMM)

- ▶ A 4MB  $\times$  32, 4-device configuration:
  - ▶ Each device has a capacity of 1Mbit  $\cdot$  8  $\cdot$  1 = 8Mbit = 1MB arranged into 8-bit words.
  - ▶ There are 4 devices; the total capacity is 4  $\cdot$  1MB = 4MB.
  - ▶ The bank, which has a 32-bit data-path, is filled by merging the devices: address  $x$  refers to

$$x \mapsto \begin{cases} \text{bits } 7 \dots 0 & \text{of MEM from device 0} \\ \text{bits } 15 \dots 8 & \text{of MEM from device 1} \\ \text{bits } 23 \dots 16 & \text{of MEM from device 2} \\ \text{bits } 31 \dots 24 & \text{of MEM from device 3} \end{cases}$$

merging each 8-bit part into the word MEM.



Notes:

## Conclusions

### ► Take away points:

1. The initial goal was an  $n$ -element memory of  $w$ -bit words; the final solution is motivated by divide-and-conquer, i.e.,
  - 1.1 one or more channels, each backed by
  - 1.2 one or more physical banks, each composed from
  - 1.3 one or more devices, each composed from
  - 1.4 one or more logical banks, of
  - 1.5 one or more arrays, of
  - 1.6 many cells
2. The major complication is a large range of increasingly detailed options:
  - lots of parameters mean lots of potential trade-offs (e.g., between size, speed and power consumption),
  - need to take care of detail: there are so many cells, any minor change can have major consequences!
3. Even so, there is just one key concept: we have some cells, and however they are organised we just need to identify and use the right cells given some address.

Notes:

## Additional Reading

- *Wikipedia: Computer Memory*. URL: [http://en.wikipedia.org/wiki/Category:Computer\\_memory](http://en.wikipedia.org/wiki/Category:Computer_memory).
- D. Page. “Chapter 8: Memory and storage”. In: *A Practical Introduction to Computer Architecture*. 1st ed. Springer-Verlag, 2009.
- A.S. Tanenbaum and T. Austin. “Section 3.3.5: Memory chips”. In: *Structured Computer Organisation*. 6th ed. Prentice-Hall, 2012.
- A.S. Tanenbaum and T. Austin. “Section 3.3.6: RAMs and ROMs”. In: *Structured Computer Organisation*. 6th ed. Prentice-Hall, 2012.
- W. Stallings. “Chapter 5: Internal memory”. In: *Computer Organisation and Architecture*. 9th ed. Prentice-Hall, 2013.

Notes:

## References

- [1] [Wikipedia: Computer Memory](http://en.wikipedia.org/wiki/Category:Computer_memory). URL: [http://en.wikipedia.org/wiki/Category:Computer\\_memory](http://en.wikipedia.org/wiki/Category:Computer_memory) (see p. 71).
- [2] D. Page. “Chapter 8: Memory and storage”. In: *A Practical Introduction to Computer Architecture*. 1st ed. Springer-Verlag, 2009 (see p. 71).
- [3] W. Stallings. “Chapter 5: Internal memory”. In: *Computer Organisation and Architecture*. 9th ed. Prentice-Hall, 2013 (see p. 71).
- [4] A.S. Tanenbaum and T. Austin. “Section 3.3.5: Memory chips”. In: *Structured Computer Organisation*. 6th ed. Prentice-Hall, 2012 (see p. 71).
- [5] A.S. Tanenbaum and T. Austin. “Section 3.3.6: RAMs and ROMs”. In: *Structured Computer Organisation*. 6th ed. Prentice-Hall, 2012 (see p. 71).

Notes: