

# Intro. to Computer Architecture

Daniel Page

Department of Computer Science,  
University Of Bristol,  
Merchant Venturers Building,  
Woodland Road,  
Bristol, BS8 1UB. UK.  
([csdsp@bristol.ac.uk](mailto:csdsp@bristol.ac.uk))

January 9, 2018

Keep in mind there are *two* PDFs available (of which this is the latter):

1. a PDF of examinable material used as lecture slides, and
2. a PDF of non-examinable, extra material:
  - ▶ the associated notes page may be pre-populated with extra, written explanation of material covered in lecture(s), plus
  - ▶ anything with a “grey’ed out” header/footer represents extra material which is useful and/or interesting but out of scope (and hence not covered).

Notes:

Notes:

- ▶ **Goal:** (briefly) introduce the topic [4, Part 1] of

**finite automata** ≡ Finite State Machines (FSMs)

in two steps, namely

1. this lecture ⇒ what FSMs are ≈ the theory
2. next lecture ⇒ design and implementation of hardware FSMs ≈ the practice

Notes:

## “Automata Theory in 10 minutes” (1)

### Definition

An **alphabet** is a non-empty set of **symbols**.

Notes:

### Definition

A **string**  $X$  wrt. some alphabet  $\Sigma$  is a sequence, of finite length, whose elements are members of  $\Sigma$ , i.e.,

$$X = \langle X_0, X_1, \dots, X_{n-1} \rangle$$

for some  $n$  st.  $X_i \in \Sigma$  for  $0 \leq i < n$ ; if  $n$  is zero, we term  $X$  the **empty string** and denote it  $\epsilon$ . It can be useful, and is common to write elements in human-readable form termed a **string literal**: this basically just means writing them from right-to-left without any associated notation (e.g., brackets or commas).

### Definition

A **language** is a set of strings.

Notes:

## Definition

A (formal) **grammar** is a tuple

$$G = (N, n, T, P)$$

including

1.  $N$ , a finite set of **non-terminal symbols** that includes a **start symbol**  $n \in N$ ,
2.  $T$ , a finite set of **terminal symbols** (which is disjoint from  $N$ ), and
3.  $P$ , a finite set of **production rules** each of the form

$$P_i : (N \cup T)^* \rightarrow (N \cup T)^*.$$

Notes:

- **Finite State Machines (FSMs)** are a model of **computation**.

- ▶ **Finite State Machines (FSMs)** are a model of **computation**.
  - ▶ An FSM is an (idealised) computer  $C$ , which, at a given point in time, is in one of some finite set of states.

Notes:

- ▶ **Finite State Machines (FSMs)** are a model of **computation**.
  - ▶ An FSM is an (idealised) computer  $C$ , which, at a given point in time, is in one of some finite set of states.
  - ▶  $C$  accepts an input string wrt. some alphabet  $A$ , one symbol at a time; each symbol induces a change in state.

Notes:

► **Finite State Machines (FSMs)** are a model of **computation**.

- An FSM is an (idealised) computer  $C$ , which, at a given point in time, is in one of some finite set of states.
- $C$  accepts an input string wrt. some alphabet  $A$ , one symbol at a time; each symbol induces a change in state.
- Once the input is exhausted,  $C$  halts; depending on the state it halts in, we say either
  1.  $C$  accepts (or recognises) the input string
  2.  $C$  rejects the input string

Notes:

► **Finite State Machines (FSMs)** are a model of **computation**.

- An FSM is an (idealised) computer  $C$ , which, at a given point in time, is in one of some finite set of states.
- $C$  accepts an input string wrt. some alphabet  $A$ , one symbol at a time; each symbol induces a change in state.
- Once the input is exhausted,  $C$  halts; depending on the state it halts in, we say either
  1.  $C$  accepts (or recognises) the input string
  2.  $C$  rejects the input string

- For a language  $L$  of all possible input strings which  $C$  *could* accept, we say

$C$  accepts (or recognises)  $L \equiv L$  is the language of  $C$

and use  $L$  to classify  $C$  ...

Notes:

## "Automata Theory in 10 minutes" (4)

### Definition

Machine	Combinatorial logic	Finite automaton	Push-down automaton	Linear-bounded automaton	Turing machine
Memory		0 stacks	1 stacks	2 stacks	2 stacks
Language		regular	context free	context sensitive	recursively enumerable
Grammar		regular ( $X \rightarrow x$ or $X \rightarrow xY$ )	context free ( $X \rightarrow \gamma$ )	context sensitive ( $\alpha X \beta \rightarrow \alpha \gamma \beta$ )	unrestricted ( $\alpha \rightarrow \beta$ )
Chomsky-Schützenberger hierarchy		type-3	type-2	type-1	type-0

### Notes:

- In the description of grammars, the idea is (read from left-to-right) there are progressively less and less restrictions:  $X$  and  $Y$  both represent non-terminal symbols,  $x$  represents a terminal symbol, and  $\alpha$ ,  $\beta$  and  $\gamma$  are strings of either terminal or non-terminal symbols, all wrt. some alphabet and grammar. So, the far right-hand case has rules that are totally unrestricted, for example, whereas the far left-hand case has rules that must be of a particular form.

## "Automata Theory in 10 minutes" (5)

### Definition

A (deterministic) **Finite State Machine (FSM)** is a tuple

$$C = (S, s, A, \Sigma, \Gamma, \delta, \omega)$$

including

1.  $S$ , a finite set of **states** that includes a **start state**  $s \in S$ ,
2.  $A \subseteq S$ , a finite set of **accepting states**,
3. an **input alphabet**  $\Sigma$  and an **output alphabet**  $\Gamma$ ,
4. a **transition function**

$$\delta : S \times \Sigma \rightarrow S$$

and

5. an **output function**

$$\omega : S \rightarrow \Gamma$$

in the case of a **Moore FSM**, or

$$\omega : S \times \Sigma \rightarrow \Gamma$$

in the case of a **Mealy FSM**,

noting an **empty** input denoted  $\epsilon$  allows a transition that can *always* occur.

### Notes:

### Question

Design an FSM that decides whether a binary sequence  $X$  has an odd number of 1 elements in it.

Notes:

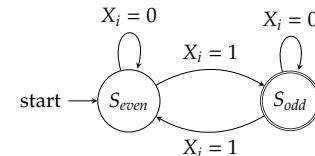
### Question

Design an FSM that decides whether a binary sequence  $X$  has an odd number of 1 elements in it.

### Algorithm (tabular)

$Q$	$\delta$	
	$X_i = 0$	$X_i = 1$
$S_{even}$	$S_{even}$	$S_{odd}$
$S_{odd}$	$S_{odd}$	$S_{even}$

### Algorithm (diagram)



Notes:

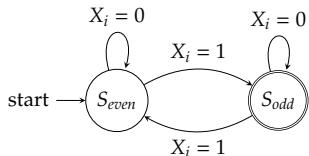
### Question

Design an FSM that decides whether a binary sequence  $X$  has an odd number of 1 elements in it.

### Algorithm (tabular)

$Q$	$\delta$
$X_i = 0$	$Q' = S_{even}$
$X_i = 1$	$Q' = S_{odd}$

### Algorithm (diagram)



### Example:

- For the input string  $X = \langle 1, 0, 1, 1 \rangle$  the transitions are

$$\rightsquigarrow S_{even} \xrightarrow{X_0=1} S_{odd} \xrightarrow{X_1=0} S_{odd} \xrightarrow{X_2=1} S_{even} \xrightarrow{X_3=1} S_{odd}$$

so the input is accepted (i.e., has an odd number of 1 elements).

- For the input string  $X = \langle 0, 1, 1, 0 \rangle$  the transitions are

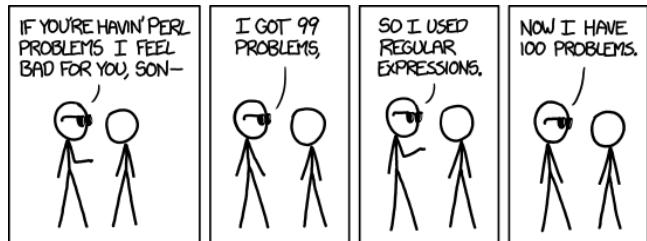
$$\rightsquigarrow S_{even} \xrightarrow{X_0=0} S_{even} \xrightarrow{X_1=1} S_{odd} \xrightarrow{X_2=1} S_{even} \xrightarrow{X_3=0} S_{even}$$

so the input is rejected (i.e., has an even number of 1 elements).

Notes:

### Some "real-world" motivation (1)

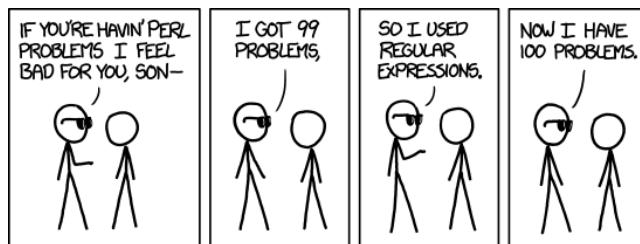
Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers



Notes:

## Some “real-world” motivation (1)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers



### ► Idea:

$$\begin{array}{lll} \text{arithmetic expression} & \xrightarrow{\text{evaluate}} & \text{number} \\ \text{regular expression} & \xrightarrow{\text{evaluate}} & \text{language} \end{array}$$

st. you can think of a regular expression (or regex) as

1. a pattern used to describe or generate a language, *or*
2. a pattern used to identify (i.e., match) members of a language.

<http://xkcd.com/1171/>

© Daniel Page (<https://csweb.cs.york.ac.uk/~dpage/>)  
Intro. to Computer Architecture

git # 3b6641 @ 2018-01-09

University of  
 BRISTOL

Notes:

## Some “real-world” motivation (2)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

### Question

We say  $X$  is a **regular expression** if it is

1. a symbol in the alphabet, i.e.,  $\{x\}$  for  $x \in \Sigma$ ,
2. the union of regular expressions  $X$  and  $Y$  st.

$$X \cup Y = \{x \mid x \in X \vee x \in Y\},$$

3. the concatenation of regular expressions  $X$  and  $Y$  st.

$$X \parallel Y = \{(x,y) \mid x \in X \wedge y \in Y\},$$

or

4. the **Kleene star** of regular expression  $X$  st.

$$X^* = \{(x_0, x_1, \dots, x_{n-1}) \mid n \geq 0, x_i \in X\}.$$

allowing for various short-hands, e.g.,

$$\begin{array}{rcl} x & \equiv & \{x\} \\ xy & \equiv & \{x\} \parallel \{y\} \\ X^+ & \equiv & X \parallel X^* \end{array}$$

Notes:

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$0^*10^* \equiv \left\{ s \mid \begin{array}{c} s \text{ is a string containing} \\ \text{a single } 1 \end{array} \right\}$$

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

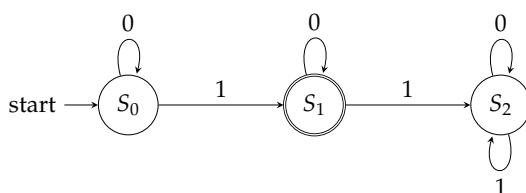
## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$0^*10^* \equiv \left\{ s \mid \begin{array}{c} s \text{ is a string containing} \\ \text{a single } 1 \end{array} \right\}$$

which can be realised using



Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$\Sigma^* 001 \Sigma^* \equiv \left\{ s \mid \begin{array}{l} s \text{ is a string containing} \\ 001 \text{ as a sub-string} \end{array} \right\}$$

which can be realised using

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

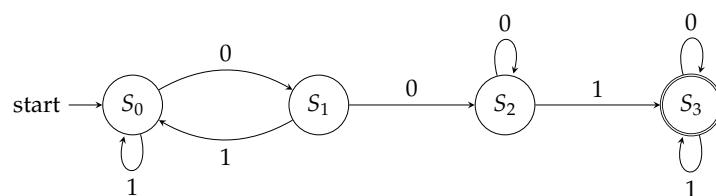
## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$\Sigma^* 001 \Sigma^* \equiv \left\{ s \mid \begin{array}{l} s \text{ is a string containing} \\ 001 \text{ as a sub-string} \end{array} \right\}$$

which can be realised using



Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$(\Sigma\Sigma)^* \equiv \left\{ s \mid \begin{array}{l} s \text{ is a string} \\ \text{of even length} \end{array} \right\}$$

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

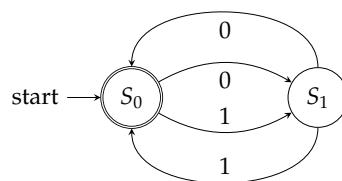
## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$(\Sigma\Sigma)^* \equiv \left\{ s \mid \begin{array}{l} s \text{ is a string} \\ \text{of even length} \end{array} \right\}$$

which can be realised using



Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$1^*(01^+)^* \equiv \left\{ s \mid \begin{array}{l} s \text{ is a string in which} \\ \text{every } 0 \text{ is followed by at least one } 1 \end{array} \right\}$$

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

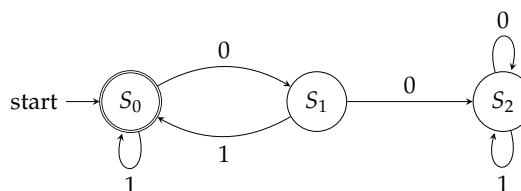
## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{0, 1\}$ , then

$$1^*(01^+)^* \equiv \left\{ s \mid \begin{array}{l} s \text{ is a string in which} \\ \text{every } 0 \text{ is followed by at least one } 1 \end{array} \right\}$$

which can be realised using



Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “*s* is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In *even* more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{a', b', \dots, z'\}$ , then

$$\text{grep -E } '.*t{o+}.*' \equiv \left\{ s \mid \begin{array}{l} s \text{ is a line read from } \text{stdin} \text{ containing} \\ \text{a 't' followed by at least one 'o' character} \end{array} \right\}$$

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “s is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In even more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

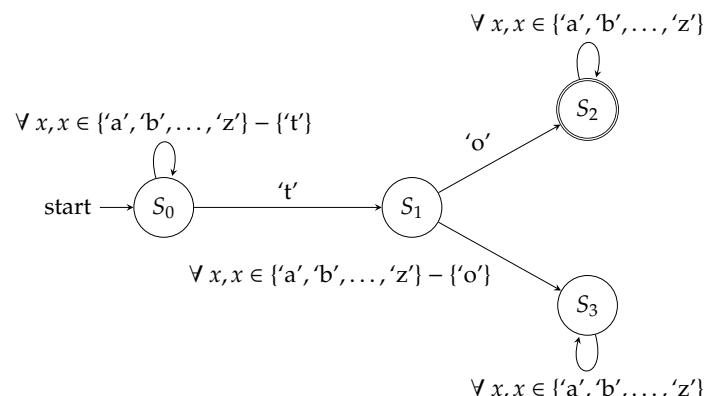
## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{a', b', \dots, z'\}$ , then

$$\text{grep -E } '.*t{o+}.*' \equiv \left\{ s \mid \begin{array}{l} s \text{ is a line read from } \text{stdin} \text{ containing} \\ \text{a 't' followed by at least one 'o' character} \end{array} \right\}$$

which can be realised using



Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “s is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In even more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{a', b', \dots, z'\}$ , then

$$\text{grep -E } '.*t\w+.*' \equiv \left\{ s \mid \begin{array}{l} s \text{ is a line read from } \text{stdin} \text{ containing} \\ \text{a 't' followed by at least one 'o' character} \end{array} \right\}$$

which can be realised using

```
1 forall lines X read from stdin do
2   Q ← s
3   for i = 0 upto n - 1 do
4     | Q ← δ(Q, Xi)
5   end
6   if Q ∈ A then
7     | print line X to stdout
8   end
9 end
```

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “s is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In even more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (3)

Example #1: regular expressions + grep  $\rightsquigarrow$  FSMs as recognisers

- Example [4, Example 1.53]: if  $\Sigma = \{a', b', \dots, z'\}$ , then

$$\text{grep -E } '.*t\w+.*' \equiv \left\{ s \mid \begin{array}{l} s \text{ is a line read from } \text{stdin} \text{ containing} \\ \text{a 't' followed by at least one 'o' character} \end{array} \right\}$$

which can be realised using

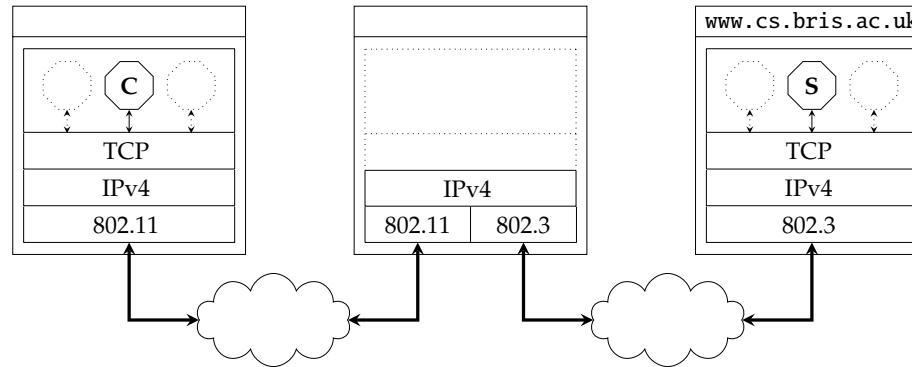
```
1 void grep() {
2   char X[ 1024 ];
3
4   while( NULL != fgets( X, 1024, stdin ) ) {
5     int n = strlen( X ), Q = start;
6
7     if( X[ n - 1 ] == '\n' ) {
8       X[ n - 1 ] = '\0'; n--;
9     }
10
11    for( int i = 0; i < n; i++ ) {
12      Q = delta[ Q ][ X[ i ] ];
13    }
14
15    if( accept[ Q ] ) {
16      fprintf( stdout, "%s\n", X );
17    }
18  }
19 }
```

Notes:

- The example implementation in the grep example is a little simplistic, so the behaviour (and hence output) will in fact differ vs. *real grep*. The technical reason for this is that we have ignored (or disallowed) backtracking: if the FSM accepts a ‘t’ followed by something other than ‘o’ then it fails, but obviously this ignores the possibility that there might be another ‘t’ later which is followed by ‘o’. So the description might be better written as “s is a line read from `stdin`, where a) there is at least one ‘t’ character, and b) the first ‘t’ character is followed by at least one ‘o’ character”. We could try to capture that using a more complex transition function, but then the example becomes overly complex, although the details is important in a general sense, the underlying idea is more important at this point. In even more detail, the difference here relates to a difference between so-called Deterministic Finite Automata (DFAs) and Non-deterministic Finite Automata (NFAs): we are dealing with and assume the former, whereas grep makes use of the latter (see, e.g., [4, Section 1.2]).

## Some “real-world” motivation (4)

Example #2: networked communication via TCP  $\rightsquigarrow$  FSMs as controllers

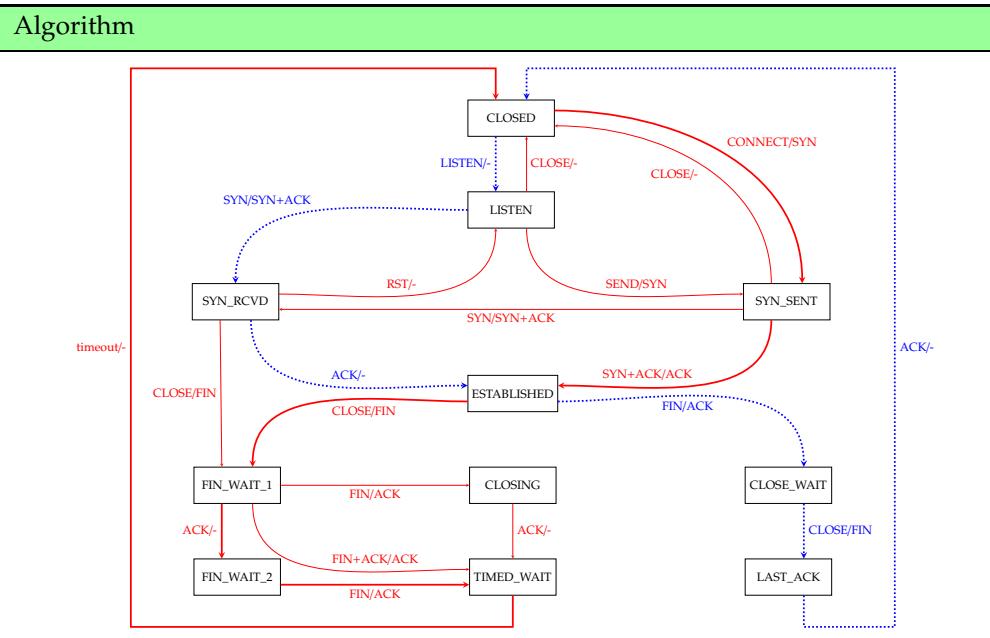


Notes:

## Some “real-world” motivation (5)

Example #2: networked communication via TCP  $\rightsquigarrow$  FSMs as controllers

### Algorithm



Notes:

## Some “real-world” motivation (6)

Example #3: typical video game “loop”  $\leadsto$  FSMs as systems



### Algorithm

```
1 reset the game state
2 while  $\neg$  game over do
3   | read control pad (e.g., check if button pressed)
4   | update game state (e.g., move player)
5   | produce graphics and/or sound
6 end
```

Notes:

## Some “real-world” motivation (6)

Example #3: typical video game “loop”  $\leadsto$  FSMs as systems



### Algorithm

```
1  $Q \leftarrow s$ 
2 while  $Q \notin A$  do
3   |  $X_i \leftarrow$  control pad
4   |  $Q \leftarrow \delta(Q, X_i)$ 
5   |  $\{\text{graphics}, \text{sound}\} \leftarrow \omega(Q)$ 
6 end
```

Notes:

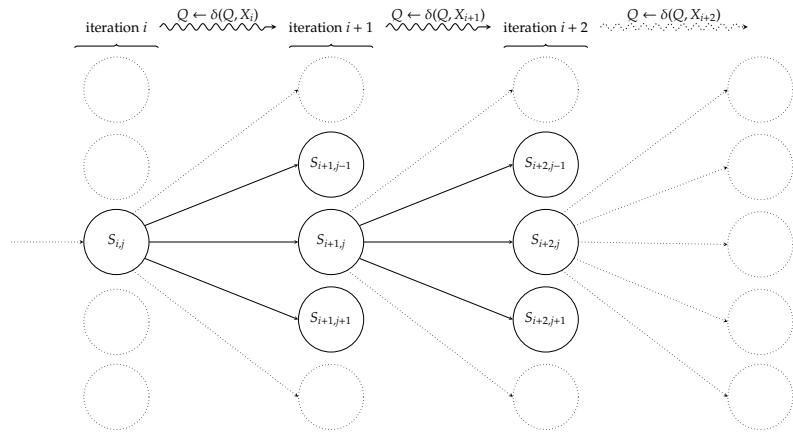
## Some “real-world” motivation (7)

Example #3: typical video game “loop”  $\leadsto$  FSMs as systems

### ► Idea:

iterations of game loop  $\leadsto$  game tree  
 $\simeq$  state space

i.e.,



which is most obvious wrt. turn-based games (e.g., chess).

Notes:

## Some “real-world” motivation (8)

### ► Take away points:

- FSMs have a variety of practical applications, e.g.,
  1. recognisers,
  2. controllers,
  3. ...
  4. *specifications*: like an algorithm, but more easily able to cater for asynchronous events
- the theory underlying which is expanded upon elsewhere.
- Even if the state space is *huge*, it may still be finite;  $\delta$  can a **partial function** st. it needn't be defined for all inputs.

Notes:

Continued in next lecture ...

Notes:

## Additional Reading

- ▶ Wikipedia: *Finite State Machine (FSM)*. URL: [http://en.wikipedia.org/wiki/Finite-state\\_machine](http://en.wikipedia.org/wiki/Finite-state_machine).
- ▶ D. Page. “Chapter 2: Basics of digital logic”. In: *A Practical Introduction to Computer Architecture*. 1st ed. Springer-Verlag, 2009.
- ▶ M. Sipster. “Chapter 1: Regular languages”. In: *Introduction to the Theory of Computation*. 2nd ed. Thomson Course Technology, 2006.

Notes:

## References

- [1] Wikipedia: *Finite State Machine (FSM)*. URL: [http://en.wikipedia.org/wiki/Finite-state\\_machine](http://en.wikipedia.org/wiki/Finite-state_machine) (see p. 75).
- [2] D. Page. “[Chapter 2: Basics of digital logic](#)”. In: *A Practical Introduction to Computer Architecture*. 1st ed. Springer-Verlag, 2009 (see p. 75).
- [3] M. Sipster. “[Chapter 1: Regular languages](#)”. In: *Introduction to the Theory of Computation*. 2nd ed. Thomson Course Technology, 2006 (see p. 75).
- [4] M. Sipster. *Introduction to the Theory of Computation*. 2nd ed. Thomson Course Technology, 2006 (see pp. 5, 37–60).

Notes: